NOÛS

NOÛS 46:4 (2012) 587–634

*Sarcasm, Pretense, and The Semantics/ Pragmatics Distinction**

ELISABETH CAMP
University of Pennsylvania

Abstract

Traditional theories of sarcasm treat it as a case of a speaker's meaning the opposite of what she says. Recently, 'expressivists' have argued that sarcasm is not a type of speaker meaning at all, but merely the expression of a dissociative attitude toward an evoked thought or perspective. I argue that we should analyze sarcasm in terms of meaning inversion, as the traditional theory does; but that we need to construe 'meaning' more broadly, to include illocutionary force and evaluative attitudes as well as propositional content. I distinguish four subclasses of sarcasm, individuated in terms of the target of inversion. Three of these classes raise serious challenges for a standard implicature analysis.

The standard view of sarcasm or verbal irony[1] was articulated by Quintilian roughly two millennia ago, as speech in which "we understand something which is the opposite of what is actually said" (95/1920, 401). Indeed, sarcasm is often presented as "the textbook case," as Robyn Carston (2002, 15) says, of the fact that speaker's meaning can come apart from sentence meaning. We also have a widely accepted pragmatic explanation of how sarcasm works. According to the standard Gricean reconstruction, in speaking

*Thanks for fertile discussion to audiences at the Conference on Meaning and Communication at the Instituto de Filosofia da Linguagem in Lisbon, the 2006 Pacific APA, the University of Rochester, the University of Western Ontario, the Workshop on Metarepresentation and Non-literal Language Use at the Centre for the Study of Mind in Nature in Oslo, and the Summer Course on Meaning, Context, Intention at the CEU University in Budapest. Special thanks to Kent Bach, David Braun, Herman Cappelen, Adam Croom, Michael Glanzberg, John Hawthorne, Larry Horn, Michael Israel, Ernie Lepore, Peter Ludlow, Paul Pietroski, Adam Sennet, David Shier, Dan Sperber, Robert Stainton, Jason Stanley, and Dmitri Tymoczko, and to an anonymous reviewer for *Noûs*.

© 2011 Wiley Periodicals, Inc.

sarcastically a speaker *A* exploits a mutually shared assumption that he could not plausibly have meant what he said. "So," Grice says,

> unless *A*'s utterance is entirely pointless, *A* must be trying to get across some other proposition than the one he purports to be putting forward. This must be some obviously related proposition; the most obviously related proposition is the contradictory of the one he purports to be putting forward (1975/1989a, 34).

Because this explanation employs the same basic explanatory tools and form of analysis as Gricean explanations of typical conversational implicatures, sarcasm seems to fit nicely into Grice's overall theoretical picture. The view that in speaking sarcastically a speaker implicates the opposite of what she actually says is so widely accepted that it rarely comes in for sustained investigation in recent debates among minimalists (e.g. Borg 2004, Cappelen and Lepore 2005), indexicalists (e.g. Stanley 2000, Szabó 2001) and contextualists (e.g. Travis 2000, Recanati 2004) about the relation between semantics, pragmatics, and 'what is said'.

Despite its storied pedigree and inherent plausibility, the standard implicature view is vulnerable to attack from two very different directions. From the left, as it were, *semanticism* argues that sarcasm is semantically encoded at the level of logical form by an operator which 'inverts' the literal meaning of the word or clause to which it applies. Meanwhile, from the right, *expressivism* denies that sarcasm or verbal irony is a matter of meaning at all, arguing instead that it serves to draw attention to a disparity between how things are and how they should be, and thereby expresses a "dissociative attitude" about some aspect of this disparity. I will argue that although these two challengers locate sarcasm at opposing ends on the spectrum of meaning, they each get something importantly right. At the same time, because they both insist that sarcasm always works in just one way, they each fail to explain the full range of data. An adequate explanation requires a more subtle and expansive understanding, not merely of sarcasm, but also of meaning more generally.

More specifically, I will defend the claim that sarcasm involves a unified operation of meaning inversion, which is manifested in distinct ways by four different subspecies of sarcasm. All four varieties *invert* something that the speaker *pretends* to mean (or presupposes someone else to have meant) relative to an evoked *normative scale*. But the target of the sarcasm, and the result of the inversion, vary widely depending on the species involved. *Propositional* sarcasm functions most like the traditional model, delivering an implicature that is the contrary of a proposition that would have been expressed by a sincere utterance. *Lexical* sarcasm delivers an inverted compositional value for a single expression or phrase. *'Like'-prefixed* sarcasm commits the speaker to the emphatic epistemic denial of a declarative utterance's focal content.

And *illocutionary* sarcasm expresses an attitude which is the opposite of one that a sincere utterance would have expressed.

Because these varieties of sarcasm vary so widely in their operative targets, their rhetorical force, and their semantic status, exclusive focus on any one in isolation from the others produces a distorted picture of sarcasm as a whole. We can only achieve a unified account of sarcasm if we adopt a more inclusive model of meaning than the traditional, truth-conditional one. On the inclusive model, meaning involves a speaker's reflexive intention to be recognized as holding some attitude; this attitude may be that of holding true a proposition, or of intending to make it true, but it may also be of some other type, such as denial or hope, or scorn. One might object that such a broad construal of 'meaning' rescues the traditional view by a terminological sleight of hand. On the contrary, I believe that the fact that sarcasm exploits and inverts such a wide range of targets brings out something we need to recognize on independent grounds: that speakers' communicative purposes and intentions often encompass more than just the exchange of information, and that illocutionary force and expressive attitudes interact in systematic ways with the presentation of truth-conditional content.

The paper is organized as follows: I offer a preliminary presentation of the semanticist and expressivist views in §1 and §2, and show in §3 that each of these extreme views suffers from irremediable failures. I present a synthetic account and identify the four subspecies of sarcasm in §4, before drawing some conclusions for the broader theory of meaning in §5.

1. Semanticism

The semantic view of sarcasm begins from a general methodological bias in favor of semantic analyses. A significant group of linguists and philosophers believe that semantics should provide the most systematic, encompassing account possible of utterances' intuitive truth-conditions. Thus, Jason Stanley (2000, 391) advocates the thesis that "all effects of extra-linguistic context on the truth-conditions of an assertion are traceable to elements in the actual syntactic structure of the sentence uttered." At the least, King and Stanley (2005, 160) write,

> Before claiming that a set of intuitions cannot be due to semantic interpretation, theorists need to have investigated all of the semantic options. For . . . claims about what can only be derived pragmatically may very well be vitiated by subsequent syntactic and semantic investigation.

Similarly, Michael Glanzberg (2005, 38) warns us that "the first moral of focus is that the appearance of being merely pragmatic can drastically deceive." To 'relegate' any linguistic phenomenon to pragmatic status at the outset, on this view, is to adopt a prematurely defeatist attitude about the scope of

semantics. Ultimately, it drives us toward "semantic skepticism," at whose extreme semantics withers away entirely, leaving only syntax and pragmatics (Chomsky 2000).

The semanticist challenger applies this general methodological bias for semantic analyses to sarcasm by noting that it appears to display three important marks of semantic status. First, it seems to be *conventional*, at least in the sense of involving a specific operation on meaning which language users learn to recognize and deploy.[2] Second, it is tightly *constrained* by sentence meaning, in a way that metaphorical meaning, for instance, is not.[3] It's true that the relation between literal and sarcastic meaning is not just one of simple negation (Fogelin 1988): the typical sarcastic meaning of an utterance of

(1) Your plan sounds *fantastic*,

is not merely that the plan is not fantastic, but that it's terrible. But the relation between the two meanings is still quite controlled: intuitively, the sarcastic meaning is the "contrary" or "opposite" of the literal one. Third, sarcasm is highly *systematic* in its application: nearly any sentence can be uttered sarcastically in some context, with results that are largely predictable without much information about the specific conversational context. Indeed, one can often identify an utterance as sarcastic and discern its core intended meaning in the absence of any contextual information, simply given its tone. In all these respects—conventionality, constraint, and systematicity—sarcasm patterns with semantic meaning, and contrasts with other standard cases of pragmatic meaning, such as Grice's (1975/1989a) letter of recommendation.

If we wanted to reflect these features in our syntax, it would be plausible to postulate a "sarcasm operator," SARC. This might be represented in surface form by an intonational contour involving heavy stress, slow rate, and nasalization (Haiman 1998, Rockwell 2000). Semantically, it could take as input a word, phrase, or sentential clause and return the most salient from among a contextually-determined set of 'contrary' items of the same syntactic type. Eminences such as Grice (1989b, 53), Bach and Harnish (1979, 33), Zwicky and Sadock (1978), and Potts (2005, 212) have all at least toyed with an analysis along these lines. For instance, Bach and Harnish (1979, 33) write,

> If . . . there is an intonational clue to the sarcastic reading, it seems that such an utterance *means* the opposite of what it means without the change of intonation, and so the speaker may well have *said* that Mac was a scoundrel (or whatever). We see no reason to deny that there are characteristic sarcastic intonation contours with semantic effects (emphasis in original).

A semantic rule along the lines of SARC isn't significantly more complex than those that have been proposed to deal with quantifier domain

restriction (Stanley and Szabó 2000), indicative and subjunctive conditionals (King and Stanley 2005), focus (Rooth 1996, Glanzberg 2005), and even metaphor (Stern 2000, Leezenberg 2001). By postulating only one new syntactic item, SARC avoids introducing significant lexical ambiguity. If we include intonation as a criterion for individuating surface forms—as we need to do already, to deal with the truth-conditional effects of focus—then SARC would not introduce additional structural ambiguity. But even if we allowed that not every sarcastic utterance has an explicit intonational marker,[4] and hence that every surface form is semantically ambiguous between sarcastic and sincere readings, the resulting theory of meaning wouldn't necessarily be less parsimonious overall than that offered by a standard pragmatic analysis. On either view, the basic interpretive steps are the same: the hearer must decide whether the utterance is sincere or sarcastic; and if the latter, determine the appropriate meaning by applying a sarcasm-specific operation of contrariety. Thus, the classic argument against positing additional semantic meanings—Grice's "Modified Occam's Razor"—gets only a weak grip here.

In addition to a general methodological bias in favor of semantic analyses and the apparent tractability of a semantic model, it also seems that we can marshal two more specific arguments for semanticism. The first is that at least in many circumstances, it is possible to report sarcastic meaning with indirect quotation, as in

(1_{IQ}) Bethany said that my plan sounds *fantastic*.

It is true that such a report is likely to be misleading or infelicitous unless it mimics the original utterance's sneering tone. But this is just what the semanticist should predict, if tone is a conventional guide to the presence of SARC at LF. Indeed, the inappropriateness of a report like (1_{IQ}) *without* a sneering tone seems to cut against the standard pragmatic analysis of sarcasm as a form of implicature: on that analysis, an unaccented utterance of (1_{IQ}) reports precisely what the original speaker did say; it's just that she said it in order to communicate something else.[5]

The motivation for employing disquotational reports as evidence for semantic status is that a primary task of semantics is to explain the intuitive truth-conditions of 'what is said', where it is assumed that ordinary reports of 'what is said' are at least a *prima facie* guide to what actually *is* said. Although considerable attention has recently been paid to the implications for semantic theorizing of felicitous indirect reports that *don't* directly echo the speaker's words, and to the fact that disquotational reports containing semantically context-sensitive terms are blocked in relevantly differing contexts (e.g. Cappelen and Lepore 1997, 2005), it is generally assumed that *if* an indirect report in a distinct context which repeats the same words as the original utterance *is* felicitous, then the intuitively reported content is semantic. And often enough, sarcasm seems to fit this pattern.

The second specific argument for semanticism is that in many cases, sarcastic interpretations can be embedded within more complex constructions. This strongly suggests that sarcasm can enter into the compositional process. And for many pro-semanticist theorists, compositionally-generated content is intimately, even definitionally, connected with semantic content. Such theorists (e.g. King and Stanley 2005, Stanley 2000, Szabó 2001) allow "weak" pragmatic effects—the "saturation" of conventionally context-sensitive expressions—to enter into composition, but deny that "strong" pragmatic effects can affect 'what is said'. Instead, they maintain that whenever the intuitive truth-conditions of a speaker's utterance differ from the truth-conditions we would expect to be delivered by semantic composition, a more sophisticated theory will reveal either that some construction is semantically context-sensitive, or else that the pragmatic effect is actually introduced after the compositional determination of 'what is said' has been completed.

Embedded sarcasm is a fairly commonplace and flexible phenomenon, as the following examples attest:[6]

(2) Since you've already made so many scintillating points this evening, I think you should let someone else voice their opinion.
(3) Because George has turned out to be such a diplomat, we've decided to transfer him to Payroll, where he'll do less damage.
(4) Because he's been such a fine friend, I've struck him off my list.
(5) If Jane is as thrilled with our plan as Bill was, then we're really in trouble.
(6) If Alice is so brilliant, then she'll be the perfect dupe for our little plan.
(7) If you come to me with one more inspired idea like that, then you're out of here.
(8) If you manage to generate one more half-baked, inconsequential idea like that, then you'll get tenure for sure.
(9) [Sun shining] If it continues to rain like this, I'll come to England more often.[7]
(10) I'm sure that the cat likes having its tail pulled.[8]

Utterances like these are not particularly strained or forced, and don't depend upon highly specific conversational contexts. I've offered a fair number to suggest that they don't exploit any single construction. Thus, we cannot simply dismiss such cases as inherently infelicitous, inappropriately artificial, or utterly unusual.

Although King and Stanley (2005) are strong proponents of a semanticist methodology, they do not endorse semanticism about sarcasm. On the contrary, they assume that in non-literal speech, speakers "knowingly express false propositions and thereby communicate true ones" (2005, 159). My claim is rather that their general arguments in favor of semanticism, combined with the specific behavior displayed by sarcasm, suggest that they themselves might be succumbing to precisely the sort of prejudicial assumption about what can and cannot count as semantic that they warn us against.

2. Expressivism

The view I call *expressivism* adopts the opposite tack from semanticism, advocating the radical exclusion of sarcasm or verbal irony[9] from the realm of meaning, much as Davidson (1978) did for metaphor. Where both the standard implicature analysis and its radical semanticist cousin treat sarcasm as a figure of speech that substitutes one propositional meaning for another, expressivist theorists like Sperber and Wilson (1981) and Clark and Gerrig (1984) argue that verbal irony "involves only one meaning: the literal one" (Sperber 1984, 130). Treating irony as a form of meaning substitution, they claim, makes it out to be a mysteriously inefficient means for communicating content that could more easily be expressed literally (Wilson 2006, 1724). Instead, we should recognize that irony is in a different line of business altogether from assertion: it draws attention to a disparity between some proposition or perspective associated with a sincere utterance of the sentence and the actual circumstances of utterance, in order to express a "dissociative attitude" toward that proposition or perspective.

This general expressivist line has been implemented in two main ways. On the one hand, Sperber and Wilson (1981, Wilson and Sperber 1992, Sperber 1994, Wilson 2006) argue that irony *echoes* or *mentions* a proposition, in order to present it as an object of ridicule. (In later versions, they relax the notion of echoing to include as potential targets not just propositions, but utterances and even moral and cultural norms.) On the other hand, Clark and Gerrig (1984; see also Kreuz and Glucksberg 1989; Walton 1990, 222–224; Kumon-Nakamura et al 1995; Recanati 2004; Currie 2006, 2010) argue that the ironic speaker *pretends* to make an assertion or other speech act, in order to mock the perspective from which it would be taken seriously. Although proponents of the echoic and pretense views have spilled considerable ink distinguishing and debating the merits of their two views, the differences are largely irrelevant for current purposes. (I'll return to the details a bit in §4.) For now, the crucial point is that both groups claim that in speaking ironically, a speaker does not undertake any genuine illocutionary act at all; rather, she mentions or pretends or 'makes as if to say' something, in order to direct attention and express an attitude.

There are two main arguments for expressivism. The first is that not just any sentence can be understood ironically in just any context, even if it is uttered with a dripping intonation. As Grice himself (1967/89b:53–4) notes, the speaker must also be interpretable as expressing an evaluative attitude:

> A and B are walking down the street, and they both see a car with a shattered window. B says, Look, that car has all its windows intact. A is baffled. B says, You didn't catch on; I was in an ironical way drawing your attention to the broken window. The absurdity of this exchange is I think to be explained by the

> fact that irony is intimately connected with the expression of a feeling, attitude, or evaluation. I cannot say something ironically unless what I say is intended to reflect a hostile or derogatory judgment or a feeling such as indignation or contempt.

The connection to a critical attitude, and specifically contempt, is bolstered by the observation that the intonational contour(s) associated with irony is closely related to the expression of negative emotion.[10] However, Wilson (2006, 1727) argues that something more specific is required than just the presence of a critical attitude:

> When some hooligan breaks my car window, I may well feel critical of him (or his behaviour). However, in normal circumstances, I could not rationally attempt to convey this feeling by saying, in an ironical tone of voice, *Look, my car has all its windows intact*.

Wilson suggests that the speaker also must be echoing, in order to criticize, some previous thought. For instance, Grice's ironic statement would become felicitous if uttered in light of a previous exchange about vandalism, where one party claimed that cars in the neighborhood are generally in excellent condition.

Wilson, and expressivists more generally, are right to draw attention to the crucial role that both evaluative attitudes and evoked thoughts play in verbal irony. But this doesn't establish their core negative claim: that irony does not involve the inversion of meaning. Implicature theorists can address this first argument by adding two further clauses to their analysis: a presupposition that someone else has endorsed the content that she makes as if to say, and an implicature that the speaker evaluates this content negatively. Similarly, a semanticist might claim that a felicitous use of SARC presupposes that someone has endorsed the embedded content, and that SARC delivers a two-pronged value as its output: the truth-conditional contrary of the embedded content, and a negative evaluation of that embedded content. Many semanticists reject the inclusion of non-truth-conditional features within semantics, and so this option would not be open to them. But they could still treat the negative evaluation as a conventional implicature, as several theorists have done for slurs (e.g. Williamson 2009).

The second expressivist argument does aim to establish the negative claim that irony is not a matter of the speaker's meaning the opposite of what she says. The general point is that the target of an ironic attitude—what the irony 'operates on'—is often something at the level of the overall pragmatic effects that would be generated by a sincere utterance, rather than the uttered sentence's semantic value or its assertoric content. Insofar as irony targets something which is itself the result of the full panoply of pragmatic

interpretation, the argument goes, it cannot be treated as a case of meaning inversion in any straightforward sense. This general argument takes several more specific forms, only some of which have been developed by expressivists themselves.

First, as several theorists have noted,[11] metaphors can be ironic, as in

(11) She's the Taj Mahal.
(12) The fountain of youth is plying his charms to the little goslings.
(13) The master tailor has stitched an elegant new suit, which he plans to debut for us at the gala ball.

Thus, (11) might be used to implicate that the woman under discussion is ugly; (12) might be used to claim that a salient older professor who is overly concerned about appearing youthful is attempting to convince some incoming graduate students to study with him; and (13) might be used to describe a famous philosopher who intends to announce an implausible new view at the APA.

If we assume that metaphorical meaning is not itself semantic, as most theorists do, then accommodating these utterances within a semantic analysis of irony or sarcasm would require that SARC first operates on the conventionally-encoded value for the relevant word or phrase, and that its output is then interpreted metaphorically. But this is highly implausible. Metaphorical interpretation is a function of the particular expressions uttered, and not just of their semantic values. For instance,

(14) Tonya Harding is the bead of raw sweat in a field of dainty perspirers.

communicates something very different, or becomes uninterpretable, if we substitute 'perspiration' for 'sweat' and 'sweating people' for 'perspirers' (Stern 2000, 222). Because metaphorical interpretation requires access to the particular words uttered, the semanticist must at least grant that SARC operates on the sentence's constituent expressions in parallel with metaphorical interpretation, with the two processes then combining in some way into a unified interpretation. But there is little independent evidence that irony is normally sensitive in the same way to the particular words or phrases uttered, as opposed to their semantic values; in particular, we don't usually get cases of substitution failure for irony analogous to that for (14).[12] Further, intuitively in utterances that do combine irony and metaphor, as in (11) through (13), irony operates on the contents that are delivered by metaphorical interpretation and not vice versa.[13] This doesn't yet establish that irony is not semantic: in particular, SARC could operate on a lexically-encoded MTHAT operator of the sort postulated by Stern (2000). But adopting this

position would significantly raise the cost of semanticism about sarcasm, by broadening the range of phenomena that require a semantic analysis.[14]

Where Stern (2000) uses the order of interpretation in combined cases to argue that metaphor is semantic, Bezuidenhout (2001, 161) takes it to show that metaphor belongs to 'what is said', where this is understood in contextualist terms, as content determined via pragmatic enrichment and modulation. She then treats irony as the 'contrary' of 'what is said', in keeping with a traditional implicature model. So far as combined cases of metaphor and irony like (11) go, this is a viable view.[15] However, both semanticist and implicature accounts of irony are strongly challenged by the fact that irony can target, not just contextually enriched 'what is said', but also the *implicatures* that a sincere utterance would have generated. Indeed, often enough the implicature is the irony's *only* target. For instance, an ironic utterance of

(15) You sure know a lot.

need not take issue with the proposition that the addressee is knowledgeable; rather, the speaker's scorn may be directed exclusively at the pretended implicature that this is admirable. Similarly, an ironic utterance of

(16) The hotel room costs a thousand dollars a night. Of course, for that you get a half bottle of Australian champagne *and* your breakfast thrown in. (Bredin 1997, 7)

targets just the implicature that the room's apparently high expense is significantly offset by the half bottle of Australian champagne and breakfast; the sentence meaning is itself presented 'straight'. Likewise, in the most likely interpretation of the following exchange, with *B*'s utterance employing an artificially cheery tone,

(17) *A*: I'm sorry Aunt Louisa is such a bother.

B: Oh, she never stays for more than a month at a time, and she always confines her three cats to the upper two floors of our house.

B does genuinely assert the utterance's semantic content; only the implicature that such visits are no imposition is ironic. And in

(18) Would you mind very much if I asked you to consider cleaning up your room some time this year?

only the *manner* of speech is ironic, along with the correlative implicature that the request is supererogatory.

The fact that irony can be directed at implicatures radically compromises the theoretical attractiveness of a semanticist analysis. Including implicatures like those in (15) through (18) within the scope of semantic composition would effectively trivialize the operative notion of compositionality and radically compromise the system's simplicity and predictability. As Stanley (2005, 230) himself says in arguing against the view that deferred reference is semantic:

> When capturing a phenomenon within the semantics would result in an unconstrained semantic theory, that suggests that the phenomenon is not semantic. For example, if in order to capture a phenomenon within the semantics, one needs to exploit resources that could allow the semantic content of 'Grass is green', relative to a context, to be the proposition that snow is white, then the phenomenon is not semantic.

Similarly, treating the implicatures in (15) through (18) as part of the semantic content which is then inverted by SARC would require assigning semantic values that depart dramatically from anything that a systematic semantic theory could plausibly deliver.[16]

The same basic objection also applies to a traditional implicature account. It is obvious that an orthodox Gricean view, which ties 'what is said' closely to "the particular meanings of the elements of [the sentence uttered], their order, and their syntactic character" (Grice 1975/1989a, 87), cannot maintain that ironic utterances of sentences like (15) through (18) mean the opposite of what they say: in these cases the speaker *does* mean what she says, and the relevant 'opposite' is of a proposition that would have been implicated by saying that. But even contextualists still distinguish between 'what is said' and implicatures. In particular, they define 'what is said' as content that is part of the speaker's "direct" or "primary" speech act, and that serves as the "springboard" for "secondary," global pragmatic interpretation—that is, for the generation of implicatures of precisely the sort that would be produced by sincere utterances of (15) through (18). Thus, even for contextualists, allowing such implicatures within the scope of 'what is said' would undermine the distinction between 'what is said' and pragmatic interpretation more generally.

So far, we've found an important class of counterexamples to the view that irony means 'the opposite of what is said'. But these cases don't directly support the expressivist claims that ironic speech only draws attention to a disparity and expresses a dissociative attitude, and does not mean anything at all. A final class of examples, offered by Kumon-Nakamura et al (1995), both place more radical pressure on treating irony as any sort of meaning inversion, and provide direct support for expressivism. In these cases, the speaker's irony is directed at the entire speech act that would be undertaken

by a sincere utterance. The most decisive examples involve ironic questions, orders, and expressives. For instance, by uttering

(19) Thanks for holding the door.

the speaker doesn't thank the hearer for *not* holding the door, or for any other 'contrary' proposition. Rather, she pretends that he has held the door and deserves thanks for doing so, where this pretense then draws attention to the fact that he rudely allowed the door to shut. Similarly, by uttering

(20) How old did you say you were?

to someone acting childish, or

(21) Could I entice you to eat another small slice of pizza?

to someone who has gobbled up the bulk of the pie, the speaker doesn't sincerely ask any question at all; instead, she pretends to ask a question in order to point out the addressee's immaturity or rudeness (Kumon-Nakamura 1995, 4).

Philosophers and linguists tend to focus almost exclusively on assertion in their theorizing about language. Often enough, this restriction is not just innocuous, but positively useful. But it also makes it easy to neglect the significant role played by illocutionary force. In particular, the illocutionary force most plausibly "contrary" to assertion is denial; and it is natural to think of denial as the assertion of the original proposition's negation (Frege 1918, Geach 1965). Analogously, an exclusive focus on assertive cases of irony also makes it appear natural to treat irony as communicating the contrary of the proposition literally expressed. However, this model breaks down dramatically when applied to irony directed at illocutionary act types other than assertion. Most illocutionary acts don't have plausible 'opposites', let alone ones that can be analyzed as the same force directed at a logically related proposition. As a result, ironic speech acts other than assertion often cannot be analyzed in exclusively propositional terms. Because standard theories of speaker's meaning focus on propositional meaning ('*meaning that p*'), it is not at all clear either that the speaker does mean anything in these cases, nor in what sense what is meant could be the "opposite" of what is said.

Examples like (19) through (21) also provide strong evidence for the positive expressivist model. Intuitively, these examples do involve the speaker pretending to make a certain speech act in order to draw attention to some disparity between the actual circumstances and the circumstances in which that speech act would be appropriate, and thereby to disparage some aspect of the current situation.[17]

3. Against Expressivism

In §1, I offered an argument for semanticism, which was primarily an argument from possibility: sarcasm appears to be susceptible to a semantic analysis, and we should prefer semantic analyses where possible. In §2, I offered an argument for expressivism, which was primarily an argument from impossibility: key cases of irony or sarcasm can't plausibly be modeled as the inversion of semantic, assertive, or even propositional content, but do appear to be amenable to an expressivist treatment. Given the two arguments' respective structures, the most natural option at this point would be to incorporate the relatively simple cases that motivated both the standard implicature view and its semantic variant within the expressivist model. In this section, I argue that this cannot be done, because often enough, speakers do commit themselves by their utterances to some content other than what they literally say. In §4, I argue that we can still provide a unified and substantive analysis of sarcasm in terms of meaning inversion, if we broaden the operative notion of meaning in an independently motivated way, to include speakers' reflexive intentions to be recognized as holding certain attitudes.[18]

3.1: Narrowly-Focused Sarcasm

The fundamental problem for an expressivist account of sarcasm is something we already encountered in the positive argument for expressivism: sarcasm is often restricted in its scope to just one element within the overall speech act. In §2, I cited cases of sarcasm that target implicatures, as in (15) through (18), to show that irony can operate at the level of a pretended speech act and not just semantic or assertive content. But these cases also demonstrate something else: that sarcasm is compatible with the speaker's genuinely committing herself to the content of what she actually says. Expressivists have never explained how their model can handle this feature of these examples.

Moreover, there is another class of cases that are even more problematic for expressivism: those in which just a word or phrase is sarcastic. Expressivists have addressed some cases of this sort, such as

(22) Jones, this murderer, this thief, this crook, is indeed an honorable fellow!

where the embedded appositive phrase presents the speaker's sincere description (Sperber 1984, 133),[19] or

(23) As I reached the bank at closing time, the bank clerk helpfully shut the door in my face,

where the entire utterance except for the adjective 'helpfully' is sincere (Wilson 2006, 1736). In effect, they claim that these cases consist of two distinct utterances, one sincere and one sarcastic, woven together into a

single sentence. However, this bifurcating analysis does not cover the cases of embedded sarcasm that we discussed in §1, where the sarcasm contributes an inverted meaning to the compositional determination of a propositional content which is itself put forward with genuine force. Unlike (22) and (23), cases like

(2) Since you've already made so many scintillating points this evening, I think you should let someone else voice their opinion.
(3) Because George has turned out to be such a diplomat, we've decided to transfer him to Payroll, where he'll do less damage.
(4) Because he's been such a fine friend, I've struck him off my list.
(5) If Jane is as thrilled with our plan as Bill was, then we're really in trouble.
(6) If Alice is so brilliant, then she'll be the perfect dupe for our little plan.

become uninterpretable if the sarcastic material is simply deleted from the asserted content.[20]

None of this is to deny that something importantly echoic, allusional, or pretending is operative in these embedded cases, along with the expression of an evaluative attitude. But by itself, this is compatible with the possibility that the echoing or pretense ultimately contributes an inverted meaning to the compositional determination of a genuinely asserted content. The clear actualization of this possibility in cases of embedded sarcasm like (2) through (6) falsifies the negative expressivist claim that there is no meaning inversion.

3.2. 'Like'-Prefixed Sarcasm

The second class of cases constitute a still clearer counterexample to the negative expressivist claim. In many dialects of American English, there is a form of explicit sarcasm which prefixes the relevant sentence with 'Like' or 'As if' and employs a sneering tone, as in

(24) Like that's a good idea.

This use isn't just a crazy invention of contemporary American adolescents: it's also found in (at least) German, as *als ob* ("as if"), in Russian, as *mozhno podumat* ("It is possible to think"), and in French, as *si tu crois* ("if you think") (Haiman 1998). In many cases, like (24), sarcastic utterances with and without 'Like' feel like stylistic variants: 'Like' seems like just one more way, along with hyperbole and dripping tone, for a speaker to provide explicit cues to her sarcastic intent. In other cases, though, sarcastic utterances prefixed with 'Like' display marked, systematic differences from bare sarcasm. (For a fuller discussion, see Camp and Hawthorne 2008.) First, sarcastic 'Like' is only felicitous when combined with declarative sentences. Thus, none of

(19) through (21) can be prefixed with 'Like'. Nor can utterances containing explicit indicators of illocutionary force, such as

(25) Frankly, she's a genius. Honestly, we should hire her immediately.

even though illocutionary adverbs like 'frankly' and 'honestly' are often employed to heighten bare sarcasm.

Second, sarcastic 'Like' is syntactically restricted to the initial position of the sentence in which it occurs, and it must take scope over the entire sentence that follows: sarcastic 'Like' cannot target a single expression, or even a sentential clause within a more complex sentence. Third, typical utterances of sentences containing sarcastic 'Like' carry explicit illocutionary force, serving to commit the speaker to something very close to the 'contrary' of the bare sentence's focal content. (I'll argue in §4.4 that they commit the speaker to its denial.) These two features combine to produce marked differences in interpretation between the same sentence intended sarcastically with and without 'Like'. For instance, the most natural interpretation of

(26) Your fine friend is here.

claims that the relevant individual is present but scorns his quality as a friend; but at least for most hearers, the same sentence prefixed with 'Like'

(26_L) Like your fine friend is here.

commits the speaker to denying that the person in question is present. Similarly, where the sarcasm in a bare utterance of

(27) The man who rescued this city from ruin is now planning to run for mayor.

might be restricted just to the presuppositions triggered by the descriptive material, an utterance of

(27_L) Like the man who rescued this city from ruin is now planning to run for mayor.

definitely commits the speaker to denying that a salient individual X is planning to run for mayor.

These observations show that 'Like'-prefixed sarcasm involves a quite specific form of meaning inversion. Further, the observed pattern of behavior displays precisely the sorts of robust, systematic constraints on implementation that standardly motivate semantic and/or syntactic analyses. Indeed, sarcastic 'Like' also exhibits another property, which even more strongly motivates a semantic analysis: it licenses Negative Polarity Items (NPIs).

NPIs are expressions, like 'ever', 'any', 'yet', 'lift a finger' and 'budge an inch', which are syntactically restricted to environments with certain inferential properties. Thus, NPIs occur happily in simple sentences containing the determiner 'No', but cannot occur when 'No' is replaced by 'Some':

(28) No dog has ever bothered me./No dog has any courage.
(29) # Some dog has ever bothered me./# Some dog has any courage.

There are many NPI licensers, and it's a matter of significant controversy what distinctive property they all share. One important trait displayed by many licensers is downward monotonicity (Ladusaw 1979): from the truth of sentences like those in (28), one can infer the truth of a more restrictive claim, as in

(30) No yellow dog has ever bothered me./No yellow dog has any courage.

This can't be the whole story about NPI licensing, not least because antecedents of conditionals and questions also license NPIs, but aren't downward entailing in any straightforward sense (Fauconnier 1978, Heim 1984, Progovac 1994, Zwarts 1995). Further, the relevant notion of 'entailment' seems to be closer to contextually-justified inference than to semantic entailment (Linebarger 1987, Krifka 1995, Israel 1996, von Fintel 1999). The precise characterization and explanation of these licensing environments thus remains a matter of heated dispute, although several theorists have focused on a connection with scalar implicature (e.g. Kadmon and Ladman 1993, Israel 2001, Chierchia 2004), which we'll return to in §4.4.

The remarkable fact about 'Like'-prefixed sarcasm is that it clearly and consistently licenses NPIs where bare sarcasm does not.[21] Thus,

(31) {Like/As if} I was going to give him any money.

is perfectly fine, on a par with

(32) It's not true that I was going to give him any money./I wasn't going to give him any money.

while

(33) # I was going to give him any money.

is terrible, even with a drippingly scornful intonation. This pattern generalizes quite freely: 'Like' licenses nearly all NPIs, including 'strong' NPIs like 'yet' and 'lift a finger', which require not just that their licenser be downward monotonic, but also anti-additive[22]:

(33) {Like/As if} anyone cares about her silly problems.
(34) {Like/As if} they're ever going to find the real killer.
(35) {Like/As if} those guys believe a word they say.
(36) {Like/As if} I've talked to George in weeks.
(37) {Like/As if} that relationship is going to last long.
(38) {Like/As if} James has ever lifted a finger to help anyone besides himself.
(39) {Like/As if} I give a damn if I ever hear a single word from you again.

All of these sentences are good, while even dripping utterances of their bare counterparts are bad.[23]

NPI licensing is standardly taken to provide compelling proof against wholesale semantic skepticism, because it seems so obvious that a predictive explanation of the observed syntactic behavior depends upon the inferential properties of the relevant constructions, such as downward-entailingness. As I noted above, it appears that those inferential properties do involve significant pragmatic enrichment. But it is also true that pragmatic context alone, even extremely overt pragmatic support, doesn't suffice to license NPIs: with very few exceptions, there must always be a semantically appropriate lexical trigger.[24] I'll sketch my positive account of 'Like' in §4.4. For now, my primary point is just that sentences containing sarcastic 'Like' do commit the speaker to some determinate content, which is in a clear sense the opposite of what a sincere utterance of the embedded sentence would have meant. This undermines a purely expressivist analysis, with its denial of any such meaning inversion. Further, both narrowly-focused sarcasm and sarcastic 'Like' clearly demonstrate that sarcastic meaning interacts with semantic meaning in more intimate ways than the standard implicature view can allow, and that strongly support some version of semanticism.

4. Varieties of Sarcasm

4.1. Sarcasm and Verbal Irony

How should we reconcile the disparate range of phenomena we've surveyed so far? One option would be to conclude that the theorists discussed in §1 and §2 are simply talking past one another: the traditional theorist is right that sarcasm involves some sort of meaning inversion, while the expressivist is right that irony proper involves drawing attention to a disparity between how things are and how they should be and expressing a "dissociative attitude" toward some aspect of this disparity. Although I agree that we need to distinguish among types of irony and sarcasm, I believe that such a radical bifurcation is at best stipulative, and at worst obscurantist. Sarcasm and verbal irony clearly do differ in some respects. Sarcasm is usually thought to be more pointed, blatant, and negative than sophisticated cases of verbal irony such as Mercutio's wry comment on his fatal wound (*Romeo and Juliet* III.1.66–67):

(41) No, 'tis not so deep as a well, nor so wide as a church door; but 'tis enough, 'twill serve. Ask for me tomorrow, and you shall find me a grave man.

At the same time, there is also wide agreement that the two phenomena are closely related. Many people use the terms nearly interchangeably. More importantly, expressivists have taken themselves to be showing that traditional theorists like Quintillian and Grice are wrong in their analysis of a relatively unified and encompassing class of utterances, of which utterances like

(1) Your plan sounds fantastic.
(42) He's a fine friend.

are paradigmatic instances.[25] Merely imposing different labels does nothing to elucidate how the various phenomena actually work, and threatens to distract us from a close examination of their similarities and differences.

Although we should expect some vague and ragged boundaries, I think we can develop a fairly systematic and substantive topography of the overall terrain if we begin by characterizing the genus of verbal irony, and then home in on sarcasm as a potentially more restricted class. In the end, I believe an analysis of sarcasm in terms of meaning inversion can explain all of the examples discussed both by traditional implicature theorists and by expressivists—so long as we are prepared to construe 'meaning' more broadly than just semantic meaning, 'what is said', or even propositional speaker's meaning, and instead as encompassing a speaker's reflexive intention to be recognized as holding some attitude.

The best characterization of the broad genus of verbal irony, I believe, derives from Kumon-Nakamura et al's (1995) "allusional pretense" theory. Their view consists of two claims, each generalizing the core commitment of the two main expressivist rivals. First, ironic utterances are *allusive*, in the sense of "call[ing] the listener's attention to some expectation that has been violated in some way"; where this violation of expectation itself entails "a discrepancy between what is expected (what should be) and what actually is" (Kumon-Nakamura et al 1995, 5). Typically, a speaker draws attention to this discrepancy in order to communicate a negative evaluation of the actual circumstances; but as Kumon-Nakamura et al note, the expressed attitude may also be positive.[26] Second, ironic utterances involve *pretense*, in the sense that the speech act is presented as not being straightforwardly genuine or sincere.[27]

I think this view is basically correct, as far as it goes. The fundamental problem with expressivism lies in its negative claim that irony *doesn't* involve meaning inversion. To their credit, Kumon-Nakamura et al remain fairly neutral about this negative claim; however, they also remain studiously silent about how meaning inversion, if it were to exist, might work. I believe that if

we adopt a more inclusive notion of meaning, then the negative claim applies at most to those cases of verbal irony that are not instances of sarcasm, and that we should identify sarcasm with the large (and possibly exhaustive) subclass of verbal irony that does involve meaning inversion.

I also think that we can exploit Kumon-Nakamura et al's model for more than just a characterization of verbal irony. If we modify each of its two major components—*allusion* and *pretense*—it can also serve as the basis for a workable analysis of sarcasm. Sarcasm, I claim, is speech which presupposes a *normative scale*; which *pretends* to undertake (or at least, evokes) a commitment with respect to this scale; and which thereby communicates an *inversion* of this pretended (evoked) commitment. I'll discuss each of these factors, and how they grow out of Kumon-Nakamura et al's model, in turn.

Consider first their notion of allusion. In effect, allusion is a form of conspicuous presupposition: presenting an assumption as if it belongs to the conversational common ground, and thereby raising it to prominence in the current context. The use of conspicuous presupposition lends irony and sarcasm much of their rhetorical force: rather than explicitly endorsing or criticizing the alluded-to assumption, or even explicitly claiming *that* someone—themselves, the hearer, or some third party—endorses it, the speaker simply acts as if that party does endorse it, and thereby makes it part of the common ground that it is so endorsed, unless the hearer actively refuses to accommodate.

In the most canonical cases of sarcasm and irony, such as

(1) Your plan sounds fantastic,

the speaker merely *pretends* to make an assertion or other speech act, but she thereby genuinely *presupposes* some standard of evaluation, and also *implicates* that this standard has been violated and that she feels negatively about its violation. Further, she accomplishes all of this without genuinely undertaking any direct assertion, question, or imperative at all.[28] In effect, in these cases the speaker attempts to manipulate the common ground without making a move that is itself recorded on the conversational scoreboard. If the hearer does acknowledge both that a presupposed expectation has been violated, and the legitimacy of the speaker's displeasure at its violation, then the perspective which underwrites this negative evaluation gains tacit acceptance without the speaker's ever explicitly articulating or defending it. The combined package of pretense, presupposition, and implicature thus carries significant rhetorical advantages for the speaker, when it works. (This is perhaps one reason that sarcasm is so frequently deployed in high school and other contexts of intense social anxiety.) However, as we'll see in §4.2, it also constitutes a kind of communicative bluff, which is particularly risky

when the relationship between speaker and hearer is antagonistic, as is often the case in sarcastic exchanges.

So far, I've just spelled out and motivated Kumon-Nakamura et al's appeal to allusion in a little more detail, as a form of conspicuous presupposition. Kumon-Nakamura et al present allusion as a single criterion. But for our purposes, it is more useful to separate it into two distinct factors: appealing to an established expectation, and drawing attention to the violation of this expectation. This provides the basis for two of our criteria. First, in place of Kumon-Nakamura et al's general appeal to an "expectation," I propose that sarcasm always presupposes or at least evokes a *normative scale*, according to which some quality, person, fact, or situation X is valorized, and others comparatively disvalued, in some ranked order.[29] Second, in place of their criterion of drawing attention to a violation of expectation, I propose that sarcasm always involves some sort of *inversion* of the evoked scale. In the canonical case, the speaker pretends to treat some situation, person, or feature Y as falling at the top of the scale, and thereby communicates that Y lies at or near the bottom. In cases of positive sarcasm, though, the situation is reversed. In either case, it is precisely because Y so flagrantly departs from the normative value exemplified by X, either positive or negative, that it warrants being assigned the contrary value from what the speaker pretends to express.

The third modification to Kumon-Nakamura et al's view centers on the role of pretense or insincerity. As I noted in §2, one of their paper's great virtues was to bring attention to non-assertional irony, which cannot be analyzed propositionally. In establishing this general point, they also cited ironic assertions, such as

(18) You sure know a lot.

where the speaker is genuinely committed to the literally encoded content. However, they failed to note the full significance of these sorts of examples: that the insincerity associated with sarcasm can be very precisely targeted. Different types of sarcasm take different 'scopes', and thereby produce very different illocutionary and rhetorical results.

Putting these modifications to Kumon-Nakamura et al's view together with a modified version of the traditional claim about meaning inversion, we get the view that sarcastic utterances presuppose a *normative scale*; they *pretend* to undertake (or at least, *evoke*) one commitment with respect to this scale; and they thereby communicate some sort of *inversion* of this pretended commitment. The sarcastic pretense can take four different 'scopes'. *Propositional sarcasm* targets and inverts a proposition that would have been associated with a sincere assertion of the uttered sentence; *lexical sarcasm* targets just a single expression or phrase within the uttered sentence; *'Like'-prefixed sarcasm* targets the focal content of an embedded declarative sentence; and

illocutionary sarcasm encompasses the entire speech act that would have been undertaken by a sincere utterance. I discuss each class in turn.

4.2. Propositional Sarcasm

The most straightforward cases of sarcasm are those in which the sarcasm's scope is directed toward some proposition to which a sincere utterance would have committed the speaker. In the simplest cases, such as

(42) He's a fine friend.
(43) James must be a real hit with the ladies.

the speaker pretends to assert the proposition *P* that is fixed by semantic composition plus lexically-focused pragmatic processes—roughly, what contextualists identify as 'what is said'. This proposition *P* evokes a situation at one extreme of an evaluative scale, typically the positive end; by pretending to assert *P*, the speaker implicates the contrary of *P*, *Q*, along with a correlative evaluative attitude toward *Q*.

So far, this conforms fairly closely to Grice's implicature model, on which the speaker means the opposite of what she says and expresses a negative attitude toward the content of what she means. But as we saw in §2, the proposition *P* which is the target of the sarcasm need not be 'what is said', understood in strict Gricean terms. First, as in (43), it may be fixed by way of pragmatic enrichment or modulation. Further, it may be determined through metaphor, as in

(11) She's the Taj Mahal,

where the speaker's metaphor evokes a scale of beauty and he pretends to claim that the woman referred to lies at its top, but in fact implicates that she belongs at or near the bottom. Propositional sarcasm can also target propositions that do not belong within contextualist 'what is said'. It can target just an utterance's presuppositions, including presuppositions generated by speech acts other than assertion:

(27_Q) Is the man who rescued this city from ruin really planning to run for mayor now?

On the most natural sarcastic reading of (27_Q), the speaker implicates that the man in question did *not* rescue the city in question from ruin, although some person or group takes him to have done so and admires him for it. Finally, as we saw in §2, propositional sarcasm can target implicatures that would be generated by a fully sincere utterance of a sentence which is itself genuinely asserted:

(17) *A*: I'm sorry Aunt Louisa is such a bother.
B: Oh, she never stays for more than a month at a time, and she always confines her three cats to the upper two floors of our house.

In (17), *B* evokes a scale of ease and burdensomeness of guests and pretends to implicate that Aunt Louisa's described actions place her at the easy end of the scale, but thereby implicates the contrary proposition: that her visits are quite burdensome.

Although these cases make trouble for a traditional view of sarcasm as meaning the opposite of what is said, they do not undermine the more general claim that sarcasm implicates the contrary of some proposition associated with the utterance. In every case of propositional sarcasm, the targeted proposition *P* is associated in some, perhaps merely pragmatic way, with some evaluative scale; and the speaker implicates the contrary of *P* with respect to that scale, along with a correlative evaluative attitude toward *Q*.

It might seem paradoxical that a speaker could implicate *Q* without first making a genuine assertion or other direct illocutionary act. In particular, I said in §4.1 that sarcastic speakers often attempt to manipulate the common ground without making a move that is itself recorded on the conversational scoreboard. But many theorists take a speaker's meaning *P* to be equivalent to her recording her commitment to *P* on the scoreboard (Lewis 1979, Brandom 1983). Thus, perhaps either the speaker of a sarcastic utterance like

(1) Your plan sounds fantastic.

must really be *asserting* that the plan is terrible, as Bach and Harnish (1979, 74) claim; or else she must not really *mean* that it is terrible at all, as expressivists maintain.

I think it is important to explaining the nuanced rhetorical role of these sarcastic utterances that the speaker merely implicates, and does not assert or otherwise directly commit herself to, the inverted content *Q*. Assertion, I take it, is an act of open and overt illocutionary commitment: a matter of placing oneself on the conversational record as committed to a certain attitude or content. But a speaker employing propositional sarcasm does not directly and overtly commit herself to *Q* in this way; this is what makes her utterance a kind of communicative bluff. As a result, a flat-footed hearer can call this bluff by responding as if the speaker really did mean *P*. For instance, a hearer *H* might respond to a sarcastic utterance of (1) with something like:

(44) Since you're so enthusiastic, let's have you present the plan to the Dean at next week's meeting.

Faced with such a response, the original speaker *S* must either acquiesce by accepting the onerous assignment, or else disavow her earlier utterance by saying something like, "I didn't really mean *P*; I was being sarcastic. I really think *Q*." A response like (44) is certainly uncharitable, in the sense that it deliberately ignores *S*'s interpretive intentions—intentions which are just as accessible as those that drive the recognition of other implicatures, and which the intentionally flat-footed hearer chooses to ignore precisely because he does recognize them. But a lack of charity does not constitute a violation of the Cooperative Principle. As we might put it, it is simply an insistence on "working to rule," not a breaking of the conversational contract.

The sarcastic speaker of an utterance like (1) makes herself vulnerable to this lack of charity precisely because she both says something she doesn't mean, and means something she doesn't say. In this respect, sarcastic utterances like (1) have a rhetorical structure similar to that of Grice's (1975/1989a) letter of recommendation for a job teaching philosophy:

(45) To whom it may concern: John's handwriting is excellent and his attendance at departmental events is punctual. Yours, etc.

In both cases, the speaker makes her communicative intention manifest to an adequately perceptive hearer, but in a way that avoids explicit commitment to the communicated content *Q*. And in both cases, one important motivation for avoiding an explicit commitment is the desire to preserve *deniability*. Thus, such speakers can legitimately object to later reports of them as having asserted or claimed *Q*—although it might be fair to report them in more general terms as having "indicated" their belief in *Q*. The difference between (1) and (45) is that the writer of (45) does mean what he actually says, while the speaker of (1) does not; and it is precisely this difference that forces the speaker of (1), faced with an aggressively uncharitable response like (44), to disavow the content of what she actually said in a way that the writer of (45) need not.

Deniability is a valuable commodity in political discourse and other contexts involving antagonistic interlocutors. Irony and sarcasm are particularly rhetorically useful because they enable speakers to communicate rhetorically volatile negative attitudes while preserving deniability (cf. Sperber and Wilson 1986, Winner et al 1988). And notably, deniability persists even when the speaker's sarcastic intent is overtly manifested in a dripping intonation; indeed, Rockwell (2000, 485) claims that we should expect reliance on nonverbal cues like tone in precisely those situations where speakers want to communicate negative messages while protecting themselves from their potentially negative consequences. When the speaker's assertion itself is merely pretended, however, as in (1), then deniability brings considerable risk.

Sarcasm contrasts sharply here with metaphor, which does have the force of assertion or another primary illocutionary act. Here too, the fact that the

speaker does not mean what she says makes her vulnerable to a flat-footedly literal response (Camp 2006a). And here too, the fact that the speaker does not explicitly say what she means allows her some leeway about the precise content of her claim. However, in contrast both to sarcasm and to utterances which merely juxtapose the two subjects that a metaphor 'yokes' together, the speaker of a metaphor cannot deny that she has committed herself to some content or other by her utterance—*even* if her main perlocutionary aim is to induce a non-propositional perspective on the subject in the hearer (Camp 2008). Thus, while expressivists about sarcasm and metaphor are right to emphasize that both of these figures of speech engage in something other than just the communication of propositions, they seriously misconstrue the communicative role that both sorts of figurative utterances actually play.[30] Further, because both expressivists and implicature theorists neglect the fine-grained rhetorical and conversational consequences of sarcasm and metaphor, they miss out on important differences between them.

What, then, about the other side of the coin: if speakers don't assert the inverted contents in sarcastic utterances like (1) and (41) through (43), why should we think that they *mean* those contents at all? Bach and Harnish (1979, 101), for instance, hold that "devious acts" of "innuendo, deliberate ambiguity, and sneaky presupposition" do not count as cases of speaker's meaning, precisely because these acts are designed to preserve deniability. The only way to achieve deniability, they think, is for the speaker to intend that her intention to get the hearer to think *Q* not itself be recognized by the hearer *as* her communicative intention; and if that is right, then such "devious acts" lack the reflexive structure essential to speaker's meaning (cf. also Lepore and Stone, 2010). I agree that some sorts of "devious" communicative acts, such as subliminal advertising, lack reflexive intentions. But innuendo, hints and insinuations are not like this: in these cases, the speaker does intend for the hearer to recognize that she is trying to communicate some attitude; and she intends for his recognition of her intention to play a crucial role in his making sense of her utterance. That is, she does intend for her communicative intentions to be manifest to the hearer. It's just that she doesn't intend those intentions to be fully publically manifest, in a way that would allow even someone who was not fully attuned to the specific, nuanced presuppositions operative in that particular conversation to recover her meaning.

By trading on this gap between what is mutually believed by the immediate conversationalists and what would be mutually believed by a wider, less informed audience, the speaker preserves at least the technical rhetorical right to pretend that she was not assuming the truth of presuppositions that are in fact crucial to deriving her actually intended meaning. The degree of deniability available to the speaker depends on how specific the operative assumptions are to the particular conversation at hand, and how highly explicit or salient they are within that conversation. At the same time, in canonical cases of sarcasm where the speaker neither means what she says

nor says what she means, the hearer can retaliate rhetorically by insisting on a narrowly literal construal of what the speaker did actually say.

4.3. Lexical Sarcasm

If propositional sarcasm most closely approximates to the traditional implicature model, lexical sarcasm provides the best case for a semanticist postulation of SARC. In cases of lexical sarcasm, as in

(3) Because George has turned out to be such a diplomat, we've decided to transfer him to Payroll, where he'll do less damage.

the speaker undertakes an overall speech act whose illocutionary force is guided by the uttered sentence's grammatical mood in the usual way, and whose content is a compositional function of the standard meanings of its constituent terms plus local, lexically-focused pragmatic processes. The notable feature, of course, is that the operative 'local processes' include inverting the meaning of at least one expression.

Lexical sarcasm displays an even tighter connection to an evoked evaluative scale than propositional sarcasm. Where the evaluative scale in propositional sarcasm might be merely pragmatically evoked, lexical sarcasm most naturally targets expressions which denote the extreme end of a conventionally-associated, normatively-loaded scale—expressions like 'brilliant', 'inspired', 'genius', 'diplomat', and 'thrilled'—so that the sarcastic inversion contributes a value at the scale's extreme other end. Often, the targeted expression denotes a positive value, but it can also be negative, as in

(8) If you manage to generate one more half-baked, inconsequential idea like that, then you'll get tenure for sure.

However, the mere presence of an evaluative scale is not sufficient to make lexical sarcasm felicitous: for instance, an utterance of

(46) If David is a real genius, then he won't get better than a C in organic chemistry.

sounds bizarre in the absence of a specific supporting context, even with sneering emphasis on 'real genius'. The additional requirement, as our model suggests, is an allusion to some other evaluation of the subject under discussion. Indeed, many if not all cases of lexical sarcasm employ explicitly allusive or comparative expressions, such as 'so', 'such a', or 'like that'—and indeed, (46) becomes significantly better if 'real genius' is replaced with 'such a genius'. These allusive expressions anchor the targeted expression's pretended evaluation to some genuine evaluation, in one of two ways. On the

one hand, they may allude to some mutually recognized and highly salient features(s) in the world which obviously warrant the opposite evaluation from the one the speaker pretends to express, as in (3). On the other hand, they may allude to a previous, genuine evaluation of the same subject with the opposite valence from the one the speaker now pretends to express, as in (8), or

(6) If Alice is so brilliant, then she'll be the perfect dupe for our little plan,

which is felicitous only if another conversational participant has just described Alice's various intellectual blunders, or if Alice's great ineptitude is otherwise extremely salient.[31]

The speaker of an utterance containing lexical sarcasm does undertake a genuine primary assertion or other illocutionary act: for instance, an utterance of

(47) Get your witty, sophisticated friends out of here now, before they cause any more damage.

does order the hearer to remove his friends. Further, the inverted normative value is not itself merely implicated; at least when located in an appropriate syntactic position, it contributes compositionally to the content of what is asserted, ordered, or asked.[32] This both considerably lessens the speaker's deniability, and makes flat-footedly literal responses along the lines of (44) more difficult, because there is often no coherent alternative assertion or other primary speech act which the speaker could pretend to have been intending to make, or which the hearer could insist on construing the speaker as having undertaken.

Although in principle propositional and lexical sarcasm differ significantly, in practice they can be difficult to tell apart, as in:

(1) Your plan sounds fantastic.
(42) He's a fine friend.
(48) That's a brilliant idea.
(49) That's a good idea.
(50) I'm sure Jane will be thrilled to hear your good news.

In some cases there may be no determinate fact of the matter about which type of sarcasm the speaker intended. But often, we can discern some differences between the two types. First, lexical sarcasm tends to employ an intonational contour emphasizing the targeted evaluative expression, while propositional sarcasm more nearly approaches an exaggerated version of a normal contour. (So, for instance, sarcastic utterances of (1) and (48) are

likely to heavily emphasize 'fantastic' and 'brilliant', while (49) will emphasize 'That'.) Second, lexical sarcasm is more likely to employ an expression at the extreme end of an evoked scale: thus, the converses of adjectives like 'fantastic' and 'brilliant' are 'terrible' and 'idiotic'—just the sorts of qualities that clearly merit sarcastic scorn. By contrast, the converse of 'good' in (49) is just 'not good', which is compatible with mere mediocrity, not itself an obviously contemptible quality. As a result, (49) is more likely to be a case of propositional sarcasm, in which the speaker implicates that it is emphatically not true that the idea is a good one, and that anyone who might think it was even good warrants scorn.[33]

It would be very tidy if lexical sarcasm were also the only type of sarcasm that embeds. And clear examples of embedded propositional sarcasm are often difficult to construct. However, they do seem to be possible, as in Levinson's example:

(9) [Sun shining] If it continues to rain like this, I'll come to England more often.

Although one could insist that the sarcasm in (9) is restricted to 'rain', or perhaps to 'this', this seems implausible; rather, it seems that the entire sentence in the antecedent is inverted. Further, propositional sarcasm embeds quite freely within epistemic modals, such as (50) and

(10) I'm sure that the cat likes having its tail pulled.

So long as the speaker's communicative intentions are sufficiently obvious, and so long as the speaker's overall utterance presents a coherent set of evaluative attitudes, it seems that embedded sarcasm can pick up on and target any aspect of the utterance.

4.4. 'Like'-prefixed Sarcasm

Like propositional sarcasm, 'Like'-prefixed sarcasm targets an entire proposition. But where bare propositional sarcasm is quite flexible in which of the various propositions associated with an utterance it can target—focal content, presuppositions, or implicatures, as generated by sentences of any grammatical mood—'Like'-prefixed sarcasm only combines with declarative sentences, and only targets content that is determined by composition of the constituent expressions' conventional meanings plus lexically-focused pragmatic processes. This inevitably includes the sentence's focal content, and often only that content.[34] Further, where bare propositional sarcasm generates at most a strong implicature that the speaker is committed to the inverted content, utterances prefixed with sarcastic 'Like' actively commit the speaker to denying that content, in a way that robustly undermines deniability. Thus,

flat-footedly sincere replies along the lines of (44) are ruled out, and speakers cannot object to reports of them as having committed themselves to the denial of that content. For instance, a speaker who uttered

(37) Like I've talked to George in weeks.

cannot pretend to have intended to claim that she has in fact spoken to George recently, and could fairly be reported as having denied speaking with George recently; while even a speaker who employed a heavily sneering tone to utter

(51) Oh, I talk with George all the time.

could pretend, albeit disingenuously, that her utterance was sincere, and could technically object to such an indirect report.[35]

Given that a speaker employing 'Like'-prefixed sarcasm does undertake a primary illocutionary act and not just an implicature, and given that sarcastic 'Like' only combines with declarative sentences, it might seem natural to analyze sarcastic 'Like' and 'As if' as forms of sentential negation, perhaps as simply elliptical for 'It's not like/as if'. This would support the intuition that in many cases, such as

(24) Like that's a good idea,

the insertion of 'Like' appears to function as a stylistic variant on bare sarcasm, and suggests that we might be able to subsume 'Like'-prefixed sarcasm within propositional sarcasm. However, in Camp and Hawthorne 2008, we argue that the distinctive behavior of 'Like'-prefixed sarcasm is better captured by analyzing 'Like' in illocutionary terms. More specifically, we propose that 'Like' encodes a function from propositions to a force/content complex with denial as its illocutionary act type.[36] This model explains most of the distinctive syntactic constraints displayed by 'Like': its restriction to declarative sentences; its prohibition from the consequents of conditionals; and its incompatibility with illocutionary adverbs like 'luckily' or 'unfortunately'. It also offers a better explanation for the infelicity of replying to 'Like'-prefixed utterances with 'That's true/false', and of substituting 'Like'-prefixed sentences for 'It's not true that/It's not like P' in response to previous utterances. Finally, it can explain the infelicity of reporting speaker's beliefs with demonstratives which are anaphoric on 'Like'-prefixed sentences, as in

(52) Like Alan has any money. # She believes that.

On our account, such reports are ruled out because the demonstrative in the second sentence has as its referent an illocutionary act rather than a

propositional content, but the latter is required to provide an appropriate object for the predicated belief.

'Like'-prefixed sarcasm is undeniably a variety of sarcasm, which involves the inversion of meaning in a strong sense of the term. However, it is less obvious that 'Like'-prefixed sarcasm exhibits the other two features I have proposed: the presupposition of a normative scale, and the pretense of undertaking one commitment in order to communicate its inverse. We should expect the conventionalization of sarcasm into an explicit marker to bring significant changes; but I think we can still discern some version of these features in operation.

We shouldn't expect 'Like'-prefixed sarcastic utterances to be pretended, or "pragmatically insincere," in the way that other sarcastic utterances are, given their explicit conventional marker and concomitant lack of deniability. However, a hint of pretense or insincerity remains in the fact that the speaker merely denies, rather than actively asserts the negation of, the targeted content's negation. Further, it is notable that cross-linguistically, the lexical items which encode this sarcastic operator always have an independent function either of echoing someone else's utterance or thought, or of presenting a situation as counterfactual. Thus, sarcastic 'Like' is closely connected to the quasi-quotative use of 'Like', as in

(54) She was like, you are so totally embarrassing me right now.

which often mimics performative elements of a mentioned utterance in addition to reporting its propositional content, and which is frequently used to present unspoken thoughts.[37] Similarly, American English employs 'As if' as an alternative to sarcastic 'Like', while German employs *als ob* ("as if"), Russian has *mozhno podumat* ("It is possible to think"), and French uses *si tu crois* ("if you think"). In their non-sarcastic applications, these expressions all serve to allude to or evoke a set of circumstances or an epistemic attitude as presupposed but not actual or actually warranted. Thus, all of these expressions at least have some significant connection to evoking a counterfactual alternative, if not to pretense *per se*.[38]

The connection to a presupposed normative scale might seem harder to discern, and its absence more damning for a unified analysis of sarcasm. In particular, where bare sarcasm frequently and easily employs intensifiers like 'very', which push the targeted content toward the extreme end of an evoked scale, and where lexical sarcasm most naturally targets expressions at the extreme end of a scale, 'Like' often combines only uneasily with such intensifiers. Thus,

(24) Like that's a good idea.

sounds considerably better than

(48_L) Like that's a brilliant idea.

even though the bare sarcastic counterpart of (48_L) (that is, [48]) would be at least as if not more natural than the bare counterpart of (24) (that is, [49]). Worse, the contents of many 'Like'-prefixed sentences, such as

(55) Like she's coming to your party.

lack any obvious connection to an evaluative scale at all.

To uncover the role that evaluative scales play in 'Like'-prefixed sarcasm, and to see more clearly the role of evocation or allusion, it is useful to return to the most surprising feature of sarcastic 'Like': its licensing NPIs. A variety of theorists have argued that polarity items are intimately connected to scales, and that NPI licensers in particular are "scale reversing" contexts. More specifically, although the details of their accounts differ considerably, Kadmon and Landman (1993), Israel (1996, 2001) and Chierchia (2004, 2009) have all argued that NPI licensing involves 'emphatic exhaustion': the emphatic presentation of content in a way that rules out all of its alternatives along some scale. Canonical 'minimizing' NPIs, like 'any' and 'a wink', denote a maximally minimal quantity. As a result, in positive contexts they make extremely weak, irrelevantly trivial statements. But once they are embedded in 'scale reversing' contexts like negation, they produce extremely strong statements, by ruling out the possibility that any alternatives higher on the scale might obtain.[39]

By themselves, these observations don't get us very far in the analysis of 'Like': even if NPIs do evoke scales, the evoked scale need not be a normative one; and many sentences embedded under 'Like' don't contain NPIs at all. However, I want to suggest that 'Like' itself serves as an operator of 'emphatic exhaustion', and so that it expresses a strong evaluative attitude which locates the embedded content at the extreme end of a scale, just as our account of sarcasm predicts. More specifically, the speaker's use of 'Like' in (55) doesn't merely deny the content expressed by the embedded sentence: it expresses the speaker's evaluation of that content as falling at the very low end of the scale of epistemic probability, and thereby 'exhausts' or rules out the assignment of any higher epistemic value to it. Further, the emphatic nature of the speaker's evaluation expresses scorn toward any higher assignment of probability to that content.[40]

Finally, the scale of epistemic probability is clearly *presupposed*, in a way that further supports our illocutionary analysis in terms of denial. An utterance of (55), like denials generally, only makes sense against the background of a presupposition that someone—whether the addressee or some other salient party—endorses at least the probability, if not the actuality, of the described situation. Indeed, the role of presupposition in 'Like'-prefixed

sarcasm is considerably more robust and constrained than in bare propositional sarcasm. For instance,

(56) Nice cool day today, huh?

could be uttered sarcastically as the start to a conversation in a context where it is manifest that the weather is uncomfortably hot and sticky. But its 'Like'-prefixed analogue,

(57) Like it's a nice cool day today.

is infelicitous as an initial remark, even if the embedded content is clearly and relevantly false. (57) does become felicitous, though, if uttered in light of a previous claim to the contrary, or of a claim that directly entails its contrary—say, that the weather has been pleasant for the past month. This contrast between the behavior of bare propositional and 'Like'-prefixed sarcasm strongly suggests that where bare propositional sarcasm merely requires some general evocation of an established normative value, 'Like'-prefixed sarcasm presupposes the specific epistemic endorsement of a particular propositional content. This presupposed endorsement constitutes a positive evaluation on the scale of epistemic probability; 'Like' then inverts this endorsement by expressing scorn toward the possibility of assigning the targeted content anywhere but the bottom of the epistemic scale.

This analysis of 'Like' as an emphatic expression of minimal epistemic probability also solves the conundrum that

(24) Like that's a good idea.

sounds considerably better than

(48_{L}) Like that's a brilliant idea.

If 'Like' simply inverted a scale evoked by the embedded sentence, then (48_{L}) should be a more emphatic statement than (24), and so more pointed and effective as a sarcastic remark. As I argued in §4.3, this is precisely how lexical sarcasm operates. But it is not how 'Like' works: instead, 'Like' denies a maximally wide range of possibilities by assigning the embedded content the lowest possible epistemic probability. And since most ideas fall far short of brilliance, even the most emphatic denial that an idea is brilliant doesn't deny much. By contrast, denying that an idea even achieves the minimal positive value of being *good* thereby also denies the possibility that it might have any greater merit. Our analysis thus predicts that, as an emphatic expression of denial, and not an expression of either internal or external negation, 'Like'

should combine more effectively with the weaker than the stronger positive evaluation.

4.5. Illocutionary Sarcasm

The final species of sarcasm to consider is what I am calling illocutionary sarcasm. These are cases, like those offered by Kumon-Nakamura et al, where the sarcasm's scope encompasses not just some element within the uttered sentence, or some proposition associated with the utterance, but the entire illocutionary act that a sincere utterance of the relevant sentence would have undertaken. Often, as we saw, this targets speech acts with an illocutionary force other than assertion:

(19) Thanks for holding the door.
(20) How old did you say you were?
(21) Could I entice you to eat another small slice of pizza?

It can also include the full range of implicatures, including especially implicatures that express evaluative attitudes such as pity, admiration, or surprise.

These are the cases for which a pretense account seems most apt: the speaker 'makes as if' to undertake a certain speech act S, where S would be appropriate in some counterfactual situation X that contrasts with the current situation Y. They are also the cases that seem least amenable to an analysis in terms of meaning inversion. Nonetheless, I think that normative scales and meaning inversion still play an important role in all of the examples discussed.

Examples like (19) through (21) function, just as expressivists claim, to evoke or allude to a situation X in which their sincere utterance would be apt; and they thereby draw attention to a disparity between X and the actual circumstances Y. But that disparity always has a specific structure: the two situations occupy opposite extremes of an evoked scale, and the speaker's utterance draws attention to the fact that Y lies at the opposite or inverse end from X. As a result, these utterances serve to express an evaluative attitude toward the actual circumstances Y which is the opposite of the attitude that they pretend to express toward X. For instance, in (19) the speaker pretends to undertake an utterance which would be appropriate if the addressee had held the door, where door-holding ranks high on a scale of politeness. This pretense draws attention to the disparity between that evoked situation and the actual one, and thereby communicates the speaker's evaluation of the addressee's actual behavior as rude. Likewise, in (21), the speaker pretends to ask a question which would be appropriate in a situation where the addressee was behaving maturely for his age, and by drawing attention to the disparity between this situation and the actual one, expresses her evaluation of the addressee's behavior as immature.[41]

A similarly encompassing pretense occurs in

(15) You sure know a lot,

with the crucial complication that here the speaker does genuinely assert the sentence's conventionally-encoded content.[42] The pretense is restricted to the implicature that the utterance is a compliment, where its complimentary status depends on an associated scale of personal virtues. Here, unlike in the previous species of sarcasm, where the speaker did genuinely endorse the evaluative scale which her utterance evoked, the speaker merely *pretends* to presuppose that knowledge ranks high on this scale. Because her utterance is sarcastic, she actually expresses skepticism toward this presupposition. This thereby implicates that other virtues (such as politeness, practicality, interest in doing more 'normal' things like hanging out) actually rank more highly, which in turn implicates that the addressee is foolish for showing off an ability that doesn't really matter. The end result is that the pretended compliment is inverted into an insult.

The same fundamental dynamic also operates in a case of positive sarcasm like

(57) Poor you, lying on the beach sipping daiquiris, without even any grading to distract you from the endless tumbling of the waves.

As in (15), the speaker of (57) does genuinely commit herself to the uttered sentence's compositionally-determined content—in this case, to the claim that the addressee is (or will be) on an oceanside vacation. Her insincerity is directed toward the evaluative attitude expressed by the initial apostrophe: that the addressee is to be pitied for being in those circumstances. This pretended negative evaluation evokes a scale of activities from onerous to enjoyable, which the speaker pretends is inverted, so that grading is treated as a great pleasure and sipping daiquiris a terrible chore. This pretense teasingly implicates her own genuine envy of the addressee's doing something that really ranks high in enjoyment.[43]

By focusing on pretense about the expression of an attitude which relies upon an evoked evaluative scale, I believe we can treat all cases of illocutionary sarcasm in terms of meaning inversion. The insight that illocutionary acts other than assertion lack well-defined opposites but can be used sarcastically undermines the traditional model of sarcasm as inverting propositional content. But neither this point, nor the essential role played by evaluative attitudes, necessarily rules out a model in terms of meaning inversion, so long as we are willing to understand 'meaning' in broader, but still fundamentally Gricean terms: as a speaker's reflexive intention to be recognized by her hearer, on the basis of her utterance, as holding some attitude, which may be partly or entirely evaluative or emotional rather than purely

truth-conditional. Speakers regularly undertake such commitments in ordinary conversation, by exploiting both conventional and rational interpretive mechanisms in all the usual ways: for instance, by employing 'thick' and thin moral terms and slurs, and by speaking metaphorically. An adequate theory of meaning, and not just of sarcasm, needs to explain this.

5. Implications and Conclusions

Although sarcasm has received considerably less attention from philosophers of language than metaphor, it arguably presents more interesting challenges for them. (Metaphor may hold more important lessons for philosophers of mind, however). The traditional implicature model, its upstart semanticist cousin, and its expressivist adversary all assume an overly narrow and simplified model of meaning and communication, which prevents them from capturing the nuanced interactions that sarcasm exploits between conventional meaning and speakers' manipulation of it, and between the determination of truth-conditions and the expression of attitudes. I hope to have developed a more adequate, unified, and substantive account of sarcasm. On my view, sarcasm always evokes a normative scale, always pretends to undertake one commitment or express one attitude with respect to this scale (or, in the case of 'Like'-prefixed sarcasm, alludes to a previous commitment), and always thereby communicates a commitment or attitude which counts as the inverse of this pretended commitment or attitude relative to that scale.

In this section, I want to articulate three main lessons from our investigation of sarcasm for the broader theory of meaning. The first major lesson is that neither semantics nor the general theory of meaning can concern themselves solely with the determination of propositional, truth-conditional content—or rather, that if they insist on so restricting themselves, then they must be supplemented with parallel, systematic and intimately interacting theories of attitude expression and illocutionary force. The semanticist, the implicature theorist, and the expressivist all analyze 'what is said', and often also 'what is meant', in narrowly truth-conditional terms. But sarcasm, like many other uses of language, intertwines the communication of information with the expression of, and exhortation to, evaluative attitudes. The exclusion of norms and emotions from the realm of meaning becomes increasingly unpalatable as the range of cases widens to include not just figurative speech like sarcasm and metaphor, but also sincere, direct, and literal speech, as in the case of slurs and moral terms.

Sarcastic 'Like' and illocutionary sarcasm introduce the complexity of illocutionary force into this mix. Both 'Like''s robust syntactic constraints and its distinctive rhetorical effects suggest that it conventionally expresses a speech act type of denial, distinct from the assertion of negation. Many theorists, most prominently Frege (1918) and Geach (1965), have held that treating denial as something other than the assertion of negation is at best

"a futile complication," and perhaps the road to logical perdition.[44] While Frege and Geach are certainly right that there are formidable challenges to assimilating denial and other speech acts within the scope of logical inference, it is remarkable that sarcastic 'Like' cannot occur within the scope of conditionals. More generally, there is growing awareness of other complex linguistic structures, such as appositive phrases, epithets and slurs, and discourse particles, which display robust 'wide-scope' logical behavior and which often concern illocutionary force and attitude expression instead of (or in addition to) truth conditions (e.g. Green 2000, Potts 2005, Williamson 2009, Siegel 2002, McCready 2008).

The second major lesson provided by sarcasm is that semantics cannot be sharply encapsulated from pragmatics, even by allowing massively more weak pragmatic 'intrusion' than we might have expected.[45] King and Stanley (2005) argue valiantly that all apparent cases of strong intrusion into the compositional determination of the primary speech act can be traced back to conventionally-encoded context-sensitivity. This view has some hope of succeeding when it comes to scalar implicatures and other generalized implicatures associated with the use of specific expressions, such as 'and':

(59) If he pulled the switch and the bomb detonated, then he's responsible for the deaths; but if the bomb detonated and he pulled the switch, then he's blameless.

But the sheer number and variety of types of embedded material makes this a daunting task, to say the least. Various scholars have noticed that metaphor also embeds quite easily:

(60) If you appoint a little Chomsky, all the sociolinguists will resign.[46]
(61) If music be the food of love, play on.[47]

So does loose use:

(62) If they send me another raw steak, I'm going to ask to speak with the manager.

and deferred reference (Nunberg 2002):

(63) If the ham sandwich pulls a runner, Bill can chase him down.

We can now add sarcasm to this list.

King and Stanley offer no explanation of embedded metaphor, sarcasm, loose use, or deferred reference, and summarily dismiss Levinson's examples of embedded 'manner implicatures' as "straightforwardly unconvincing"

(2005, 153, fn. 53). However, the ordinary speakers I've consulted all find the examples I've cited to be perfectly acceptable.

Philosophers of language with a semanticist orientation, like King, Stanley, and Szabó, thus face a stark choice. On the one hand, they can hold on to the view that all of these apparent cases of strong intrusion are really generated at LF as part of the conventionally-driven compositional process. This is at least somewhat plausible for the case of lexical sarcasm, given the constraint and systematicity of sarcastic inversion that we used to motivate a semanticist analysis in §1—although the attractiveness of this model is significantly tempered by the fact that not all cases of sarcasm, or even embedded sarcasm, are lexical. There is also some (though I believe, much less) plausibility to modeling metaphor semantically, as Stern (2000) and Leezenberg (2001) have done. But loose use and especially deferred reference display very little systematicity. As Stanley (2005) acknowledges, postulating LF representations for all of these possible uses would effectively trivialize the notion of conventional composition and undermine the search for predictive power that motivated a methodological bias for semantic analyses in the first place.

On the other hand, semanticists can allow strong intrusion into the compositional process. Grice himself eventually chose the latter option, at least for generalized implicatures:

> It certainly does not seem reasonable to subscribe to an absolute ban on the possibility that an embedding locution may govern the standard nonconventional implicatum rather than the conventional import of the embedded sentence (1989c, 375).

However, I think it's fair to say that the range of cases in which an embedding locution governs something other than conventionally-encoded meaning is much wider than the 'standard' nonconventional implicata that Grice envisioned.

The admission of strong intrusion into composition might seem like a clear triumph for contextualists, since most theorists who have drawn attention to strong intrusion have done so in order to argue for a permissive, contextualist notion of 'what is said' (e.g. Travis 2000, Bezuidenhout 2001, Recanati 2003). However, even the most permissive contextualists want 'what is said' to be intuitively tied to the modulation or enrichment of lexical meaning, and regularly claim that only the enriched or modulated meaning has any psychological reality.[48] But sarcasm does not intuitively belong within 'what is said' understood in this way: on the contrary, it's a commonplace, and not a rarefied theoretical postulate, that a sarcastic speaker means the opposite of what they say. Indeed, those who argue for including metaphor within 'what is said' or semantics often do so by contrasting metaphor with irony (e.g. Bezuidenhout 2001, Stern 2000). In my (2006a), I argue that ordinary intuitions and

practices actually support excluding metaphor from 'what is said' as well. But lexical sarcasm demonstrates even more clearly that non-encoded processes can help to determine the content of the speaker's primary assertion or other illocutionary act without intuitively belonging to 'what is said'. Thus, the relationship between conventionally-encoded, compositional meaning and the speaker's primary speech act—let alone her overall communicated content—is considerably more complex than contextualists typically allow.

The idea that speakers can use particular expressions to mean something radically different from their conventional meaning is not new, of course. Among others, Clark (1983) discusses "nonce" word use; Crimmins (1998) claims that "Fregean" uses of names depend upon "semantic pretense"; and Hills (1997, 147) argues that we need to acknowledge metaphorical meanings as "full participants in the familiar recursive rigmarole of compositional semantics." In a related vein, Kripke (1977) and Davidson (1986) discuss unintentional divergences between a word's conventional meaning and what a speaker uses it to mean, occasioned by a speaker's misinformation either about the world or about her words. However, these sorts of cases have received surprisingly little attention in recent debates about semantics, pragmatics, and 'what is said'.

To avoid bogging down in a terminological quagmire, I suggest we abandon the battle over how to interpret 'what is said', which is ambiguous in ordinary speech between a locutionary use—roughly equivalent to 'what the uttered sentence means'—and an illocutionary one, equivalent to 'what the speaker claimed' (Camp 2007). Instead, I propose that we distinguish four distinct classes of commitments that individuals undertake in making their utterances. All deserve to be called 'meaning', and all potentially include evaluative and emotional attitudes as well as truth-conditional beliefs and intentions:

- '*what is locuted*'[49]: the result of composing conventionally-encoded meaning, including disambiguation and assigning values to conventionally context-sensitive terms; also including the assignment of an illocutionary-act-type correlative to grammatical mood, but without entailing actual illocutionary commitment. Roughly equivalent to 'what is said' in a strict Gricean sense.
- '*what is asserted/asked/ordered*': the speaker's primary illocutionary act, subject to local, non-encoded modifications to the compositional contributions of expressions, including metaphor, lexical sarcasm, loose use, and deferred reference. Roughly equivalent to 'what is said' in a contextualist sense.
- '*what is (nonconventionally) implicated*': a speaker's further illocutionary commitments, which are fixed via global, rational interpretive processes that take as their input the speaker's having (pretended to) assert, ask, or suggest what she did, plus the Cooperative Principle, Maxims, and assumptions specific to the particular conversational context.
- '*what is perlocuted*': further effects, including especially directing attention and non-propositional perspectival effects, which the speaker intends to produce

> in the hearer by way of his recognizing her reflexive communicative intentions. These go beyond the speaker's illocutionary act in requiring for their success that the effect actually be produced in the hearer, as opposed to the hearer's mere "uptake" of the speaker's communicative intention.

Speakers and hearers regularly display their implicit sensitivity to all four levels of meaning in the course of ordinary conversation. In particular, the conversational status of sarcasm and metaphor underscores the important distinction between 'what is locuted' and 'what is asserted'. On the one hand, in both sorts of figurative speech, ordinary speakers and hearers recognize both the authority of 'what is locuted' when faced with uncharitable, non-accommodating responses, and also the deniability of what is not actually locuted. On the other hand, both lexical sarcasm and metaphor can contribute to determining 'what is asserted/asked/ordered'.

The other three species of sarcasm also exploit these levels of meaning in systematic ways. 'Like'-prefixed sarcasm denies the content of what would be asserted by a sincere utterance of the embedded declarative sentence. Propositional sarcasm implicates the negation of some proposition associated with what is asserted/asked/ordered or implicated. And illocutionary sarcasm implicates the contrary evaluative attitude from one that would be expressed by a sincere utterance. If we insist on lumping all of these kinds of meaning together, we cannot clearly describe the similarities and differences among these types of sarcasm, or among other types of utterances more generally.

The third and final lesson is in a way the inverse of the first two. Although an investigation of sarcasm helps to bring out just how intimate and nuanced the interactions are among propositional contents, evaluative attitudes, and illocutionary force, and between conventional meaning and contextual exploitation, it also reveals that these interactions depend upon complex, systematic structures, with genuine constraints on all sides. Not just anything goes. In a deep sense, sarcasm is something that speakers *do* with their utterances. But this general operation can take a variety of fairly precise 'scopes', with distinctive, robust consequences in each case, and with at least one—'Like'-prefixed sarcasm—staking a strong claim to semantic status.

In this respect, sarcasm is not so different from negation.[50] Negation, too, can be accomplished by various means, some semantic and others pragmatic, each with distinct effects: from external, propositional negation ('It's not the case that') to internal, lexical negation ('non'-) and illocutionary denial ('I deny/reject the claim that'), and on to merely implicit rejection ('Yeah, right'/'Anyway, as I was saying . . .'). Sarcasm is considerably more complex, in that it also involves evaluative attitudes, and engages in pretense. But in both cases, we can discern a unified phenomenon across its various manifestations; and in both cases we can make substantive claims about the behavior of each of the various species. Thus, what we might have expected

to be a case study in the erosion of conventional, compositional meaning in the face of multivalent pragmatic manipulation instead demonstrates that semantics does not always wither away upon closer inspection into syntax plus pragmatics. Semanticism about sarcasm may wildly overreach, but semantic skepticism—along with its more temperate cousin, expressivism—advocates a radical retreat that is just as unwarranted.

Notes

[1] There is significant diversity of opinion about the relation between sarcasm and verbal irony: some people use the terms interchangeably or only use the latter term; many treat irony as the more general phenomenon; and some even treat sarcasm as more general. My own intuition comports with Quintilian's: sarcasm is typically more explicit than irony, and involves a simpler mapping from literal to figurative meaning. My main target here is a theory of sarcasm. In §2 I shift to talk of "verbal irony," because this is the term employed by the theorists I discuss there. In §4 I discuss the relation between the two phenomena, arguing that my analysis of sarcasm can accommodate most if not all of the cases described as verbal irony.

[2] Sarcasm appears not to be conventional in the sense of being merely arbitrary. That is, sarcasm appears to be a robustly cross-cultural phenomenon, although when and how people speak sarcastically is highly culture-dependent.

[3] Stanley (2005, 229) suggests that such constraint is the "mark of the semantic."

[4] People do seem to associate a distinctive tone with sarcasm—roughly, slow rate, exaggeratedly modulated stress, and nasalization—in the sense that this is the tone they typically adopt if explicitly asked to utter a sentence "sarcastically," or if they read out a passage one of whose sentences is clearly intended sarcastically (Haiman 1998, Rockwell 2000). However, there is also evidence that this tone is simply an instantiation of a more general expression of negative emotion. Nor does it appear to be the tone that speakers always employ when they use sarcasm in the course of normal conversation: Bryant and Fox Tree (2005, 2002; cf. also Attardo et al 2003) found that "dripping" sarcasm is higher-pitched, is 'flatter' (involves less amplitude variability), and exhibits no difference in duration in comparison to literal utterances. Finally, only some sarcasm is "dripping"; for instance, Gibbs (2000, 18) found that 24% of sarcastic utterances were uttered without any special intonation. According to Bryant and Fox Tree (2005), "dripping" sarcasm is identifiable as such using either prosodic cues alone (with lexical information eliminated through acoustic processing) or lexical information alone (in written form), while "dry" sarcasm is identifiable only given both prosodic cues and lexical information. (Attardo et al 2003 also draw attention to visual markers, such as "blank face.") Some speakers, of course, relish appearing sincere to a naïve audience, relying entirely on the violation of contextual assumptions to clue the sophisticated hearer in to their communicative intent.

[5] In my (2007), I argue that 'say'-reports are ambiguous between reporting locutionary and illocutionary acts. A locutionary interpretation of (1_{IQ}) without any special intonation would indeed be felicitous if the speaker cancelled the possibility of an illocutionary interpretation by adding a clause, as in:

($1_{IQ}'$) Bethany said that my plan sounds fantastic, but she didn't mean it—she was being sarcastic.

I think this datum needs to be taken seriously. However, many theorists take 'say'-reports to (univocally) report illocutionary acts; thus, this is the interpretation I attribute to the proponents of both semantic and pragmatic analyses of sarcasm in the text.

[6] Two classes of cases need to be set aside. First, there are utterances in which the entire sentence, or at least the consequent, is sarcastic, as in

(i) If you want a tasty, healthy, gourmet meal, then you should head over to KFC.
(ii) If he's been rejected by four women in a month, then he must be a real charmer.

Second, there may be metalinguistic constructions, such as

(iii) He's not a *real genius*, he's a genuinely good philosopher!

uttered in response to a free-standing sarcastic utterance of the antecedent. The possibility of metalinguistic negation can't show anything about semantic content *per se*, because it is "a device for objecting to a previous utterance on any grounds whatever, including the . . . implicata it potentially induces, its morphology, its style or register, or its phonetic realization" (Horn 1989, 363).

[7] Examples (4) and (9) are from in Levinson 2000, 210; (4) is originally due to Ivan Sag.

[8] From Bach and Harnish 1979, 71.

[9] As I noted in fn. 1, so far I have been talking about sarcasm, but both expressivists and traditional implicature theorists like Grice discuss verbal irony. In §4, I will argue that sarcasm is a species of verbal irony; here, I follow expressivists' talk of verbal irony, where its denotation at least largely overlaps with that of sarcasm.

[10] See fn. 4.

[11] Stern (2000, 235), Bezuidenhout (2001, 161), Camp (2006a, 291). Example (11) is from Bezuidenhout.

[12] I discuss some cases of irony directed at manner below (e.g. ex. [18]). Indeed, I think this is the best analysis of (12) and (13); see fn. 43 below. But this does not affect the point in the text, about the relative order of metaphorical and ironic interpretation.

[13] Bach and Harnish (1979, 69) claim that ironic interpretation generally precedes metaphorical interpretation. But they attempt to establish this by appealing to just one, fairly conventionalized example ('hot/cold' applied to cars), and by suggesting that the order of interpretation doesn't matter in that case. Stern (2000, 235) argues that it is only in conventionalized cases that the order of interpretation is reversible. See Popa (2009, ch. 9) for extensive argument that metaphor is interpreted before irony in combined cases, including discussion of psycholinguistic evidence.

[14] See my (2005) for criticism of Stern's view of metaphor.

[15] I argue against contextualist assimilations of metaphor to 'what is said' in my (2006a).

[16] Note in particular that the implicatures in (15) through (18) are highly particularized, unlike the scalar implicatures that theorists like Chierchia (2004) have argued are marked within the grammar.

[17] Note that in (19) through (21), the "dissociative attitude" is directed toward the actual circumstances, and not a person or perspective associated with the pretended utterance, as Sperber and Wilson maintain. The expressivist should thus merely claim that irony draws attention to a disparity between the actual circumstances and some evoked one, and expresses some evaluative attitude about this disparity. See §4.1.

[18] In this section, I revert to talking about 'sarcasm' instead of 'verbal irony'; I discuss the relation between the two terms in §4.

[19] Sperber (1984) originally cited this example to argue against a pretense view, on the ground that there is no overall speech act which a speaker could coherently pretend to undertake, while an echoic view can allow that just some constituents are mentioned or echoed. Currie (2006, 124) responded that a pretense account can permit speakers to move in and out of pretended and genuine speech, something that Wilson (2006) now concedes.

[20] Note that embedded sarcasm can contribute to quantifier domain restriction, as in (2). (Thanks to Stephen Neale for this point.) Perhaps, given that the echoic theory models irony as a form of mention akin to quotation, proponents of that view might argue that the relevant expressions in embedded sarcasm are both used and mentioned, on analogy to

(iv) I'll have you know that 'that lying S.O.B.' is my father, and I'd ask you to treat him with more respect.

The difficulty is that in cases of mixed use-mention in indirect quotation, like (iv), the speaker echoes, while possibly distancing herself from, someone else's preferred mode of presentation, but while still preserving the intended reference. In cases of embedded sarcasm like (2) through (6), by contrast, the relevant expressions contribute a radically distinct property from their conventionally-encoded referent. Thus, in (3) George is claimed to be a blunderbuss; and in (5), Jane (and Bill) are claimed to be highly critical of the plan.

Unlike Sperber and Wilson, Currie (2006, 113, fn. 6) does acknowledge that sarcasm genuinely embeds within complex constructions; he points toward Blackburn's and Gibbard's work on expressivism in ethics as a solution. The problem with this response is that the sarcastic expressions in examples like (2) through (6) don't merely express a brute pro- or con-attitude, akin to 'Boo!' or 'This is to-be-done'; they contribute a specific blend of description and evaluation, which is related to the expression's conventional meaning in a determinate, predictable way. To accommodate this fact, the expressivist must be willing to offer an expressivist analysis of all 'thick' terms. Currie doesn't appear to advocate expressivism as a general analysis of evaluative terms—and of course there are serious problems for analyses like Blackburn's and Gibbard's (for recent discussion, see Schroeder 2008). But further, if Currie were willing to adopt expressivism as a general analysis, then this would undermine the central expressivist claim about sarcasm: that it accomplishes something different in kind from ordinary talk. If expressivism is warranted for sarcastic uses of expressions like 'diplomat', 'thrilled' and 'brilliant', then it should also be adopted for those expressions' conventional meanings, since the two are inverses of one another.

[21] One reader reports hearing at least some bare sarcastic utterances containing NPIs as marginally acceptable if uttered with a very heavy tone and implied reference to prior discourse; but even for them the contrast between NPIs in bare and 'Like'-prefixed sarcasm remains dramatic. Linebarger (1987, 349) notes that (bare) sarcasm does not license NPIs, but does not note that 'Like' does license them. Horn (2001, 178) notes that sarcastic constructions like 'Fat chance' and 'Small thanks' can license NPIs, and Horn (2009) extends the point to 'Like' and 'As if'. Bender and Kathol (2003) also note that 'Like' licenses NPIs.

[22] In the sense of Zwarts (1995): for anti-additive functions, a disjunction in their scope is equivalent to a wide-scope conjunction, so that "Nobody ate or drank" entails "Nobody ate and nobody drank."

[23] Michael Israel (p.c.) offers some examples of NPI constructions which 'Like' does not license:

(v) #Like she woke up until 8.
(vi) #Like she can seem to figure this out.

Likewise, Larry Horn (p.c.) offers the following infelicities:

(vii) #Like/As if I've been back there again in years.
(viii) #Like I'm gonna get tenure until I write a book.

While this brings out the complexity of NPI licensing, note that none of these are licensed in the paradigmatic licensing environments of conditionals or questions, either.

[24] Horn (2001, 2009) argues that the orthodox requirement of overt licensing must be weakened, in response to what he calls "Flaubert licensing"—so-called because it is a kind of negation that, like "God in the deist universe and the author in the Flaubertian novel," is "everywhere present yet nowhere visible" (Horn 2001, 176):

(xi) San Francisco is beating *anyone* these days as often as the Atlantic City Seagulls beat the Harlem Globetrotters.

(xii) The tone [of Germaine Greer's attack on manufacturers of vaginal deodorants] wasn't light-hearted, which might have justified touching the subject *at all*.

However, the range of cases which permit Flaubert licensing are still quite restricted, and seems to require a negative or 'exhaustive' content as its primary communicative point. Further, Larry Horn (p.c.) also points out that 'retro-Not', as in

(ix) She's a good teacher. Not!

does not license NPIs:

(x) # James has ever lifted a finger to help you. Not!

Thus, the point remains that in general, even flagrant pragmatic signaling of a 'negative' or otherwise downward-entailing proposition does not suffice for NPI-licensing.

[25] For instance, although Wilson (2006, 1732) denies that (what she calls) verbal irony constitutes a natural kind, because it shades off into other phenomena such as indirect quotation, she still assumes that there is a unified phenomenon of verbal irony, which the echoic theory explains, and which includes all of the cases discussed in §§1 and 3.

[26] More generally, Gibbs (2000, 18) found that while 90% of sarcastic utterances involved mocking, only 54% were critical, and only 69% made a negative point by saying something apparently positive.

[27] Kumon-Nakamura et al claim that irony involves "pragmatic insincerity" rather than "insincerity at the substantive level," as occurs with lying. Their distinction between substantive and pragmatic forms of insincerity is misleading; the appropriate contrast is rather between the insincerity's overtness or covertness.

[28] I defend the claim that the speaker does not claim her communicated content in §4.2.

[29] The most familiar scales are semantic, in the sense that they impose this ordering by entailment; but as we'll see, scales can also be generated pragmatically, via contextually salient assumptions (cf. Horn 1989). Thus, a scale may be implicitly operative even if no expression within the embedded sentence directly entails a quantitative ordering.

[30] An expressivist analysis of irony as merely 'drawing attention to' a discrepancy between expectation and reality suffers from a further problem, which is precisely parallel to one faced by Davidson's analysis of metaphor as merely 'drawing attention to' a similarity between two situations. Expressivists cannot explain the fact that ironic and metaphorical utterances can be genuinely informative, in the sense of making a (potentially false) claim about something which the hearer isn't just failing to remember or notice. For further discussion on difficulties for 'juxtaposition' theories of metaphor, see my (2006b).

[31] The allusion operative in this second class differs from those we have discussed so far: instead of the speaker pretending to conform to an evoked expectation in order to draw attention to its actual violation, here the speaker pretends to express the contrary of the evoked expectation, in order to uphold it. Even here, though, I suspect that there must still be at least a hinted allusion to someone's genuinely subscribing to the evaluation that the speaker merely pretends to express: in (6), for instance, there is at least a suggestion that Alice is an ideal dupe because she believes herself to be smarter than she is.

[32] In some cases, like (47) or the examples discussed by expressivists, such as

(23) As I reached the bank at closing time, the bank clerk helpfully shut the door in my face,

the inverted material is not part of the focal, 'at issue' content. But as examples like (3) and (6) show, this is a function of syntactic position only.

[33] I return to the difference between 'good' and 'brilliant' in §4.4, in connection with 'Like'-prefixed sarcasm.

[34] Note that although sarcastic 'Like' prefers to target focal content, it can at least sometimes target presuppositions as well: for instance, to my ear the preferred reading of

(27_L) Like the man who rescued this city from ruin is now planning to run for mayor.

expresses a sincere presupposition that the man in question satisfies the description; but the sentence might also be used to express skepticism about his efficacy in addition to denying that he is running for mayor. Second, just like propositional sarcasm, sarcastic 'Like' can often target content fixed by metaphorical interpretation, as in

(11′) Like she's the Taj Mahal.
(xiii) Like the fountain of youth has managed to convert anyone to his bizarre doctrines in years.
(xiv) Like I'd ever make the beast with two backs with him.

I hope to return to the interaction between sarcasm and metaphor, and to iterated sarcasm, in future work.

[35] Note that there is a reading of sentences like (37) on which the speaker could be claiming to have spoken with George: one with a significant pause between 'Like' and the rest of the sentence. However, this reading is not available for (37), because 'introductory' 'Like', as we might call it, does not license NPIs; it is also not intersubstitutable with 'As if' in the way that sarcastic 'Like' is in (37) and more generally.

[36] For this analysis to be plausible, it is essential to acknowledge that no linguistic expression has the magical property of engendering a commitment for its utterer simply in virtue of tokening the expression—as Frege (1918) and Davidson (1979) point out, actors, journalists engaged in quotation and court recorders all token expressions without actually undertaking the concomitant illocutionary acts. However, as Mitchell Green (1997, 2000) has argued, it is possible for an expression or construction to have an illocutionary-act *type* as its semantic value, in the sense that *if* one does undertake an outright commitment to a sentence containing that expression or construction, *then* one thereby undertakes an illocutionary act of the relevant type.

[37] Though related, the two uses of 'Like' are again clearly distinct: only sarcastic 'Like' can freely intersubstitute with 'As if,' and the quasi-quotative use can be applied to non-declarative sentences.

[38] Further, Michael Israel (p.c.) notes that 'Like', 'As if' and 'Als ob' all have a literal function as comparatives, which presupposes the non-identity of the compared items, and thereby draws attention to the disparity between them.

[39] Israel (2001) discusses a range of NPIs other than the canonical minimizers, and incorporates them into a more sophisticated scalar analysis. I use them here merely to bring out the role of emphatic exhaustion, which can happen in a variety of ways.

[40] Note that bare propositional sarcasm can also target epistemic probability as its evaluative scale, as in:

(xv) {I'm sure/surely/Oh, yeah,} he's been talking to George a lot these days.

In contrast, I hear a sarcastic utterance of the embedded sentence, without any epistemic modifer, as only marginally acceptable. On my view, this is because the contents of the embedded sentence are evaluatively neutral and so lack an appropriate normative scale to target, an absence which the addition an expression of epistemic certainty remedies.

[41] Likewise, in Mercutio's wry pun in (41), he pretends that death is a positive goal which the wound will facilitate, in order to wryly implicate that it is a highly negative outcome.

[42] At least, on Kumon-Nakamura et al's imagined interpretation. We can also imagine a context in which the speaker communicates that the addressee does not possess much knowledge; this would then be a case of propositional sarcasm.

[43] I noted above, in §4.2 and fn. 34, that metaphor can fall within the scope of both propositional and 'Like'-prefixed sarcasm, as in

(11) She's the Taj Mahal.
(xiv) Like I'd ever make the beast with two backs with him.

In these cases, the speaker employs a metaphorical mode of expression to determine a content which is then inverted or denied. By contrast, the combination of metaphor and sarcasm in cases like (12) and (13) seems more properly classified as illocutionary sarcasm:

(12) The fountain of youth is plying his charms to the little goslings.
(13) The master tailor has stitched an elegant new suit, which he plans to debut for us at the gala ball.

In these cases, the speaker does genuinely assert or ask the content which would be fixed by a sincere metaphorical utterance of the relevant sentence. The sarcasm is directed exclusively at the appropriateness of this metaphorical mode of description, and at the presuppositions about the subject under discussion that would be required to support it. The same basic phenomenon is also operative in the over-polite request

(18) Would you mind very much if I asked you to consider cleaning up your room some time this year?

Manner-directed sarcasm that is not strictly metaphorical can also be used to generate content that serves as the input to the compositional determination of larger contents, as in

(xv) If Her Majesty has completed her disquisition on my many and fulsome merits, we should stop blocking traffic and get out of the intersection.

Perhaps this should be thus classified as a case of illocutionary sarcasm as well. Finally, I take it that sarcastic utterances which employ "illocutionary force indicating devices," such as

(25) Frankly, she's a genius. Honestly, we should hire her immediately,

also count as instances of illocutionary sarcasm, insofar as the sarcasm's target includes the frankness and honesty with which the speaker pretends to be putting forward her utterance.

[44] Analogously, the dominant strand of theorizing about NPI licensing, deriving from Ladusaw (1979), treats it in terms of truth-conditional entailments. But this leaves unexplained, for instance, the fact that questions license NPIs. Horn (2001, 2009) and Israel (2001) suggest that the most promising analysis of NPI licensing focuses on the utterance's illocutionary and rhetorical force rather than on its narrowly inferential properties.

[45] On the other hand, Glanzberg turns out to be right in warning that "the appearance of being merely pragmatic can drastically deceive", insofar as one variety of sarcasm—sarcastic 'Like'—is undeniably semantic.

[46] From Levinson 2000, 210.

[47] From Shakespeare, *Twelfth Night* I.i.1–3; Hills (p.c.) has used this case to show that metaphor can embed.

[48] Recanati (2003, 301, fn. 3) conjectures that all local pragmatic processes are unreflective.

[49] David Braun proposes this term in his (ms.).

[50] Jason Stanley (p.c.) suggested this analogy.

References

Attardo, Salvatore, Jodi Eisterhold, Jennifer Hay, and Isabella Poggi (2003): "Multimodal Markers of Irony and Sarcasm," *Humor* 16:2, 243–260.

Bach, Kent and Robert Harnish (1979): *Linguistic Communication and Speech Acts* (Cambridge, MA: MIT Press).

Bender, Emily, and Andreas Kathol (2003): "Constructional Effects of *Just Because . . . Doesn't Mean . . .*," *Proceedings of the 27th Annual Meeting of the Berkeley Linguistics Society* (Berkeley, CA), 13–25.

Bezuidenhout, Anne (2001): "Metaphor and What is Said: A Defense of a Direct Expression View of Metaphor," *Midwest Studies in Philosophy* 25, 156–186.

Borg, Emma (2004): *Minimal Semantics* (Oxford: Oxford University Press).

Brandom, Robert (1983): "Asserting," *Noûs* 17, 637–650.

Braun, David (ms.): "Implicating Questions."

Bredin, Hugh (1997): "The Semantic Structure of Verbal Irony," *Journal of Literary Semantics* 16:1, 1–20.

Bryant, Gregory and Jean Fox Tree (2005): "Is There an Ironic Tone of Voice?" *Language and Speech* 48:3, 257–277.

——— (2002): "Recognizing Verbal Irony in Spontaneous Speech," in *Metaphor and Symbol* 17:2, 99–117.

Camp, Elisabeth (2008): "Showing, Telling, and Seeing: Metaphor and 'Poetic' Language," *The Baltic International Yearbook of Cognition, Logic, and Communication*, vol. 3: *A Figure of Speech* [Online], ed. E. Camp, 1–24.

——— (2007): "Prudent Semantics Meets Wanton Speech Act Pluralism," in *Context-Sensitivity and Semantic Minimalism: New Essays on Semantics and Pragmatics*, ed. G. Preyer and G. Peter (Oxford: Oxford University Press, 2007), 194–213.

——— (2006a): "Contextualism, Metaphor, and What is Said," *Mind and Language* 21:3, 280–309.

——— (2006b): "Metaphor in the Mind: The Cognition of Metaphor," *Philosophy Compass* 1:2, 154–170.

——— (2005): "Critical Notice of Josef Stern's *Metaphor in Context*," *Noûs* 39:4, 715–731.

——— and John Hawthorne (2008): "Sarcastic 'Like': A Case Study in the Interface of Syntax and Semantics," *Philosophical Perspectives* 22:1: *Language and Logic*, ed. J. Hawthorne (Oxford: Wiley-Blackwell, 2008, 1–21).

Cappelen, Herman and Ernie Lepore (2005): *Insensitive Semantics* (Oxford: Blackwell Press).

——— (1997): "On an Alleged Connection Between Indirect Speech and the Theory of Meaning," *Mind and Language* 12:3/4, 278–296.

Carston, Robyn (2002): *Thoughts and Utterances* (Oxford: Blackwell).

Chierchia, Gennaro (2004): "Scalar Implicatures, Polarity Phenomena, and the Syntax/Pragmatics Interface," in *Structures and Beyond: The Cartography of Syntactic Structures* vol. 3, ed. A. Belletti (Oxford: Oxford University Press), 39–103.

Chomsky, Noam (2000): "Language as a Natural Object," in *New Horizons in the Study of Language and the Mind* (Cambridge: Cambridge University Press).

Clark, Herbert (1983): "Making Sense of Nonce Sense," in *The Process of Language Understanding*, ed. G. B. Flores d'Arcais and R. J. Jarvella (London: Wiley).

——— and Gerrig, Richard (1984): "On the Pretense Theory of Irony," *Journal of Experimental Psychology: General* 113, 121–126.

Crimmins, Mark (1998): "Hesperus and Phosphorus: Sense, Pretense, and Reference," *Philosophical Review* 107:1, 1–47.

Currie, Gregory (2010): *Narratives and Narrators: A Philosophy of Stories* (Oxford: Oxford University Press).

——— (2006): "Why Irony is Pretense," in *The Architecture of the Imagination*, ed. S. Nichols 111–133.

Davidson, Donald (1986): "A Nice Derangement of Epitaphs," in *Truth and Interpretation: Perspectives on the Philosophy of Donald Davidson*, ed. E. Lepore (Oxford: Blackwell, 433–446).

——— (1979): "Moods and Performances," in *Meaning and Use*, ed. A. Margalit (Dordrecht: Reidel), 9–20.
——— (1978): "What Metaphors Mean," *Critical Inquiry* 5:1, 31–47.
Fauconnier, Gilles (1978): "Implication Reversal in Natural Language," in *Formal Semantics and Pragmatics for Natural Languages*, ed. F. Guenthner and S. J. Schmidt (Dordrecht: D. Reidel), 289–301.
Fogelin, Robert (1988): *Figuratively Speaking* (New Haven: Yale University Press).
Frege, Gottlob (1918/1977): "Negation," in *Logical Investigations*, ed. G. Frege (Blackwell, Oxford, 31–33).
Geach, Peter (1965): "Assertion," *Philosophical Review* 74, 449–465.
Gibbs, Ray (2000): "Irony in Talk among Friends," *Metaphor and Symbolic Activity* 15, 5–27.
Glanzberg, Michael (2005): "Focus: A Case Study on the Semantics/Pragmatics Boundary," in *Semantics vs. Pragmatics*, ed. Z. Szabo (Oxford: Oxford University Press).
Green, Mitchell (2000): "Illocutionary Force and Semantic Content," *Linguistics and Philosophy* 23: 435–473.
——— (1997): "On the Autonomy of Linguistic Meaning," Mind 106:422, 217–243.
Grice, Paul (1975/1989a): "Logic and Conversation," in *Studies in the Way of Words* (Cambridge MA: Harvard University Press, 22–40).
——— (1978/1989b): "Further Notes on Logic and Conversation," reprinted in *Studies in the Way of Words*, 41–57.
——— (1989c): "Retrospective Epilogue," in *Studies in the Way of Words*, 339–385.
Haiman, John (1998): *Talk is Cheap: Sarcasm, Alienation, and the Evolution of Language* (Oxford: Oxford University Press).
Heim, Irene (1984): "A Note on Negative Polarity and Downward Entailingness," in *Proceedings of NELS* 14, ed. C. Jones and P. Sells (Amherst, MA: University of Massachusetts GLSA), 98–107.
Hills, David (1997): "Aptness And Truth In Verbal Metaphor," *Philosophical Topics* 25:1, 117–153.
Horn, Laurence (2009): "Hypernegation, Hyponegation, and *Parole* Violations," *Berkeley Linguistics Society* 35.
——— (2001): "Flaubert Triggers, Squatitive Negation, and Other Quirks of Grammar," in *Perspectives on Negation and Polarity Items*, ed. J. Hoeksema, H. Rullmann, V. Sanchez-Valencia and T. van der Wouden (Amsterdam: John Benjamins, 173–200).
——— (1989): *A Natural History of Negation* (Chicago: University of Chicago Press).
Israel, Michael (2001): "Minimizers, Maximizers and the Rhetoric of Scalar Reasoning," *Journal of Semantics* 18: 297–331.
——— (1996): "Polarity Sensitivity as Lexical Semantics," *Linguistics and Philosophy* 19:6, 619–666.
Kadmon, Nirit and Fred Landman (1993): "Any," *Linguistics and Philosophy* 4.2, 279–298.
King, Jeffrey C. and Jason Stanley (2005): "Semantics, Pragmatics, and The Role of Semantic Content," in *Semantics vs. Pragmatics*, ed. Z. Szabo (Oxford: Oxford University Press).
Kreuz, Roger and Richard Roberts (1995): "Two Cues For Verbal Irony: Hyperbole And Ironic Tone Of Voice," *Metaphor and Symbolic Activity* 10, 21–30.
Krifka, Manfred (1995): "The Semantics and Pragmatics of Polarity Items," *Linguistic Analysis* 25, 209–257.
Kripke, Saul (1977): "Speaker's Reference and Semantic Reference," *Midwest Studies in Philosophy* 2, 255–276.
Kumon-Nakamura, Sachi, Sam Glucksberg, and Mary Brown (1995): "How About Another Piece Of Pie? The Allusional Pretense Theory Of Discourse Irony," *Journal of Experimental Psychology: General* 124:1, 3–21.
Ladusaw, William (1982): "On the Notion 'Affective' in the Analysis of Negative-Polarity Items," in *Journal of Linguistic Research* 1:2, 1–16.

Leezenberg, Michiel (2001): *Contexts of Metaphor*. Current Research in the Semantics/Pragmatics Interface, vol. 7 (Oxford: Elsevier Science).
Lepore, Ernie and Matthew Stone (2010): "Against Metaphorical Meaning," *Topoi* 29:2, 165–80.
Levinson, Stephen (2000): *Presumptive Meanings: The Theory of Generalized Conversational Implicature* (Cambridge MA: MIT Press).
Lewis, David (1979): "Scorekeeping in a Language Game," *Journal of Philosophical Logic* 8, 339–359.
Linebarger, Marcia (1987): "Negative Polarity and Grammatical Representation," *Linguistics and Philosophy* 10: 325–387.
Martin, Robert (1992): "Irony and the Universe of Belief," *Lingua* 87, 77–90.
McCready, Eric (2008): "What 'Man' Does," *Linguistics and Philosophy* 31:6, 671–724.
Nunberg, Geoffrey (2002): "The Pragmatics of Deferred Interpretation," in *Blackwell Encyclopedia of Pragmatics*, eds. L. Horn and G. Ward, (Oxford: Blackwell).
Popa, Mihaela (2009): *Figuring the Code: Pragmatic Routes to the Non-literal* Dissertation, University of Geneva.
Potts, Christopher (2005): *The Logic of Conventional Implicatures* (Oxford: Oxford University Press).
Progovac, Ljiljana (1994): *Negative and Positive Polarity: A Binding Approach* (Cambridge: Cambridge University Press).
Quintilian (95/1920): *Institutio Oratoria*, trans. HE Butler, Loeb Classical Library (Cambridge MA: Harvard University Press).
Recanati, François (2004): *Literal Meaning* (Cambridge: Cambridge University Press).
——— (2003): "Embedded Implicatures." *Philosophical Perspectives* 17:1, 299–332.
Rockwell, Patricia (2000): "Lower, Slower, Louder: Vocal Clues of Sarcasm," *Journal of Psycholinguistic Research* 29:5, 483–495.
Rooth, Mats (1996): "Focus," in *Handbook of Contemporary Semantic Theory*, ed. S. Lappin (Oxford: Blackwell), 271–297.
Schroeder, Mark (2008): *Being For: Evaluating the Semantic Program of Expressivism* (Oxford: Oxford University Press).
Siegel, Muffy (2002): "Like: the Discourse Particle and Semantics," *Journal of Semantics* 19:35–71.
Sperber, Dan (1984): "Verbal Irony: Pretense or Echoic Mention?" *Journal of Experimental Psychology*: *General* 113:1, 130–136.
——— and Deirdre Wilson (1981): "Irony and the Use-Mention Distinction," in P. Cole (ed.). *Radical pragmatics* (New York: Academic Press), 295–318.
Stanley, Jason (2005): "Semantics in Context," in *Contextualism*, ed. G. Peters and G. Preyer (Oxford: Oxford University Press), 221–53.
——— (2000): "Context and Logical Form," *Linguistics and Philosophy* 23:4, 391–434.
——— and Zoltan Szabó (2000): "On Quantifier Domain Restriction," *Mind and Language* 15:2&3, 219–261.
Stern, Josef (2000): *Metaphor in Context* (Cambridge MA: MIT Press).
Szabó, Zoltan (2001): "Adjectives In Context," In *Perspectives on Semantics, Pragmatics, and Discourse*, ed. R. Harnish and I. Kenesei. (Amsterdam: John Benjamins), 119–146.
Travis, Charles (2000): *Unshadowed Thought: Representation in Thought and Language* (Cambridge, MA: Harvard University Press).
von Fintel, Kai (1999): "NPI-Licensing, Strawson-Entailment, and Context-Dependency," *Journal of Semantics* 16: 97–148.
Walton, Kendall (1990): *Mimesis as Make Believe* (Cambridge, MA: Harvard University Press).
Williamson, Timothy (2009): "Reference, Inference, and the Semantics of Pejoratives", in *The Philosophy of David Kaplan*, ed. J. Almog and P. Leonardi (Oxford: Oxford University Press), 137–158.

Wilson, Deirdre (2006): "The Pragmatics of Verbal Irony: Echo or Pretence?" *Lingua* 116, 1722–1743.

——— and Dan Sperber (1992): "On Verbal Irony," *Lingua* 87: 53–76.

Winner, Ellen, Jonathan Levy, Joan Kaplan and Elizabeth Rosenblatt (1988): "Children's Understanding of Non-Literal Language," *Journal of Aesthetic Education* 22:1, 51–63.

Zwarts, Frans (1995): "Nonveridical contexts," *Linguistic Analysis* 25:3/4, 286–312.

Zwicky, Arnold and Jerrold Sadock (1975): "Ambiguity Tests and How to Fail Them," in J. Kimball (ed.), *Syntax and Semantics* vol. 4, (New York: Academic Press, 1–36).